

COM1009

Introduction to Algorithms and Data Structures

Topic 06: Lower Bounds and Sorting in Linear Time

Essential Reading:
Chapter 8 up to page 198.

► Aims of this topic

- To discuss the class of **comparison sorts**: sorting algorithms that sort by comparing elements.
- To show a general **lower bound** for the running time of this class of sorting algorithms. They can't do better than this.
- **Other approaches can sometimes do better.** We'll see how to sort numbers in a **bounded range** in **linear time**.

► Comparison Sorts

- InsertionSort
- SelectionSort
- MergeSort
- HeapSort
- QuickSort (see Exercise Sheet 4)
- All of these by comparing elements – we call these **comparison sorts**.
- We'll see today that there's a limit to how fast they can run, but...
 - Sometimes we can go faster by using **extra information**

► Performance of Comparison Sorts

- The best comparison sorts we have seen so far take time $\Omega(n \log n)$ in the worst case.
 - InsertionSort: worst case = $\Theta(n^2)$
 - MergeSort, HeapSort: worst case = $\Theta(n \log n)$
- Can we do better?
 - Having a better algorithm could be useful
- Can we prove it's impossible to do better?
 - Stops us wasting time looking for something that doesn't exist

► Complexity Theory

(very briefly - more in COM2109 Automata, Computation and Complexity)

- Deals with the **difficulty of problems**.
- Investigates **limits to the efficiency** of algorithms
 - Results like: “no algorithm can solve problem X faster than worst case time T”.
 - Stops us from wasting time trying to achieve the impossible
 - Informs the design of efficient algorithms.
- Two sides of the same coin:

Complexity theory \leftrightarrow Efficient algorithms

► Appetiser: P and NP

(not relevant for the assessment in COM1009, but relevant for Computer Science)

- **P and NP** (much more in COM2109)
 - P = problems that can be **solved** quickly (i.e. in polynomial worst-case time)
 - NP = problems whose solutions can be **verified** quickly
- Sometimes we can **verify** solutions quickly, but we can't (currently) **find** solutions quickly.
 - Example: Given that $n > 0$ is a product of two primes, find the two primes p and q that satisfy $pq = n$
 - *Finding p and q can be hard, but if you tell me p and q , I can easily multiply them together and check whether the answer equals n*
 - Try it! See https://en.wikipedia.org/wiki/RSA_numbers

► Appetiser: NP-Completeness

(not relevant for the assessment in COM1009, but relevant for Computer Science)

- **NP-complete problems** (much more in COM2109)
 - >3000 important but seemingly different problems: satisfiability, scheduling, selecting, cutting, routing, packing, colouring, ...
 - If you can solve any **one** of these problems quickly, **then every other** problem in NP can also be solved quickly (i.e. $P = NP$)
 - But **can** any of them be solved quickly? No one (currently) knows.
- **The P vs NP Problem**
 - Is every easily-verified problem also easy to solve in the first place?
 - \$1m prize for a definite solution: <https://www.claymath.org/millennium-problems/millennium-prize-problems>

► How (Not) to Show Lower Bounds

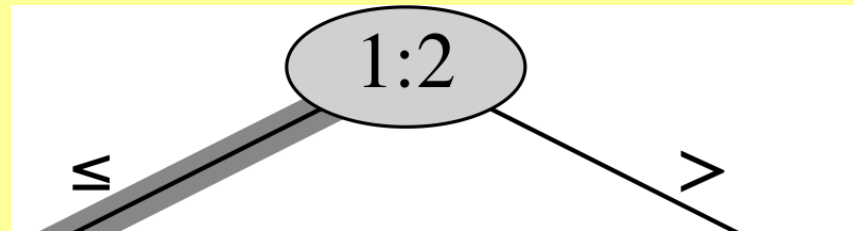
- How can we show that time $\Theta(\dots)$ is best possible?
- *“We didn’t manage to find a better algorithm.”*
 - Perhaps you didn’t look hard enough
- *“No one has ever found a better algorithm.”*
 - How do you know this? What if tomorrow someone does?
 - We have to find arguments that apply to **all algorithms that can ever be invented**.
- *“Surely, every efficient algorithm must do things this way.”*
 - Intuition is often wrong. For example, efficient algorithms for multiplying matrices start by *subtracting* elements!

► Comparison Sorts as Decision Trees

- There is one thing that all comparison sorts have to do: **compare elements**
- Let's strip away all the overhead, data movement, looping, recursing, etc. and take the number of comparisons as a lower time bound.
- We assume that elements a_1, \dots, a_n are distinct – then we can assume that all comparisons have the form $a_i \leq a_j$.
- A **decision tree** reflects all comparisons **a particular comparison sort** makes, and how the outcome of one comparison determines future comparisons.
 - Like a skeleton of a sorting algorithm.

► Decision tree for a comparison sort

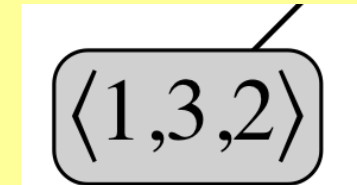
- Inner node $i:j$ means comparing a_i and a_j .



- Leaves: ordering $\pi_1, \pi_2, \dots, \pi_n$ established by the algorithm:

$$a_{\pi_1} \leq a_{\pi_2} \leq \dots \leq a_{\pi_n}$$

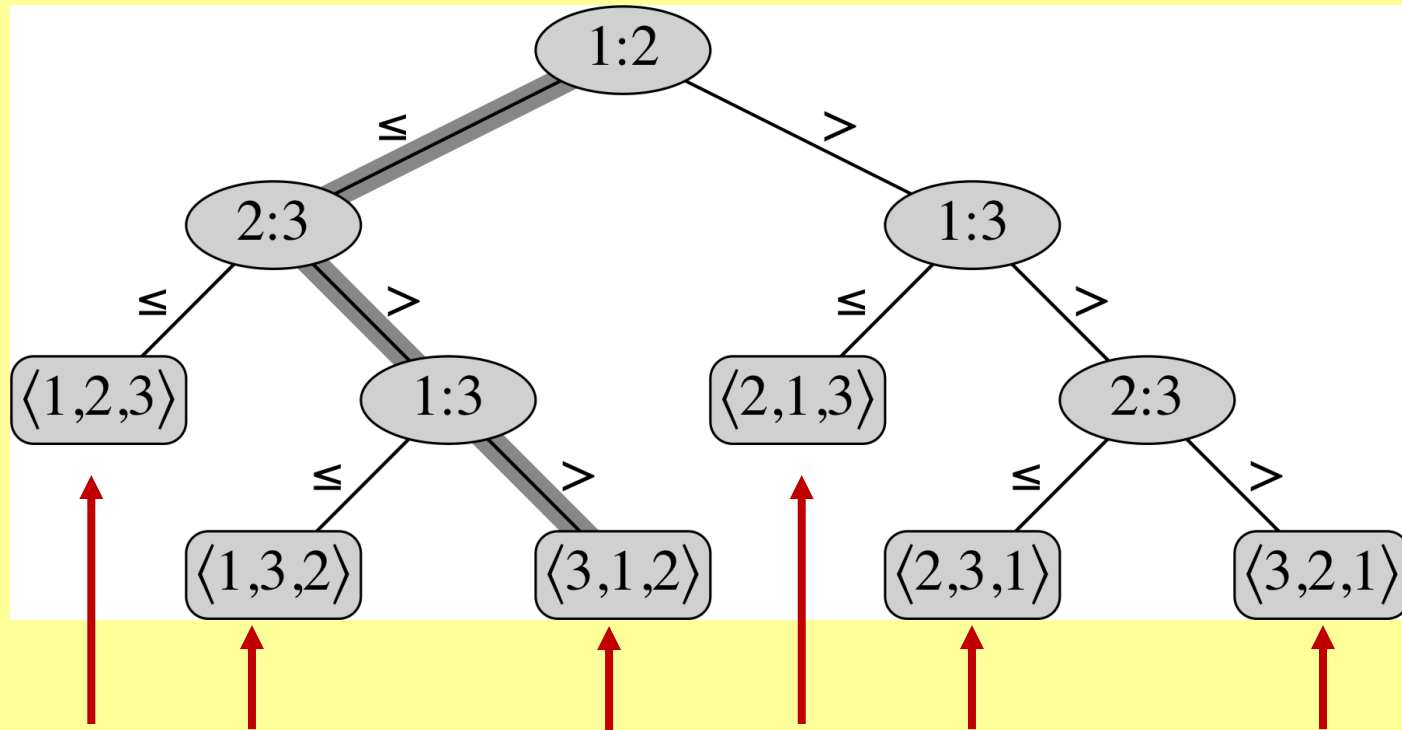
A leaf contains a sorted output for a particular input.



$$a_1 \leq a_3 \leq a_2$$

- **The execution of a sorting algorithm corresponds to tracing a simple path from the root down to a leaf.**

► Example of a decision tree



The leaves include all the possible orderings of the n values, so there must be $n!$ of them (in this example, $3! = 6$)

► Lower bound for comparison sorts

Theorem: Every comparison sort requires $\Omega(n \log n)$ comparisons in the worst case.

- This includes all comparison sorts that will ever be invented.
- Proof follows (see Theorem 8.1 in the book).
- The theorem can be extended towards an $\Omega(n \log n)$ bound for the **average-case time** (not done here).
- The theorem implies that **HeapSort** and **MergeSort** are asymptotically **optimal comparison sorts**. They achieve the best possible $\Omega(n \log n)$ worstcase runtime.

► Proof of the lower bound

- The **worst-case number of comparisons** equals the **length of the longest simple path** from the root to any reachable leaf: we call this the **height h** of the tree (as in HeapSort).
- Every correct algorithm must be able to produce a sorted output for each of the $n!$ possible orderings of the input.
- A binary tree of height h has no more than 2^h leaves.
 - We'll prove this formally in a bit; let's take this for granted for now.
- To accommodate $n!$ leaves we need $2^h \geq n!$ And so $h \geq \log(n!)$.
- So the worst-case number of comparisons is at least **$\log(n!)$**

► How big is $\log(n!)$?

- $\log(n!) = \log(n) + \log(n-1) + \dots + \log(1)$
 - So $\log(n!) \leq n \log(n)$
- $\log(n!) = \log(n) + \dots + \log(n/2) + \dots + \log(1)$
 - So $\log(n!) \geq \log(n) + \dots + \log\left(\frac{n}{2}\right)$
 - So $\log(n!) \geq \frac{n}{2} \cdot \log\left(\frac{n}{2}\right) = \frac{n}{2} \log(n) - \frac{n}{2} \log(2)$

n terms,
biggest is
 $\log(n)$

$n/2$ terms,
smallest is
 $\log(n/2)$

Lower order terms
can be ignored

$$\log(n!) = \Theta(n \log(n))$$

So worst case comparison sorting must take $\Omega(n \log n)$

► Can we do better?

- The lower bound of $\Omega(n \log n)$ is bad news for applications where comparisons are the only source of information.
- However, it suggests a way out: where possible, **use more information** than mere comparisons!
- Elements to be sorted are often **numbers or strings**, which reveal more information.
 - e.g. representing a value as a decimal (rather than a pile of pebbles) requires work; we have to represent the ones, the tens, the hundreds, ..., separately. This generates information.
 - We can use this information: is a 2-digit number bigger than a 1 digit number? We can answer this without even knowing what the numbers are.

► CountingSort: Idea

- Assume that the input elements are integers in $\{0, \dots, k\}$. This is extra information – we know up-front that the inputs are no bigger than k .
- For each element x , **CountingSort counts the number of elements less than x** .
 - For instance, if 17 elements are smaller than x , then x belongs in output position 18.
 - Beware: we need to make sure that **equal elements** are put in **different** output positions.
- CountingSort uses an array $C[0 \dots k]$ for counting and an array $B[1 \dots n]$ for writing the output.

► CountingSort

- Initialise counter array
- Count elements
- Running sum: #elements $\leq i$
- Write elements to output

COUNTINGSORT(A, B, k)

```
1: let  $C[0 \dots k]$  be a new array
2: for  $i = 0$  to  $k$  do
3:      $C[i] = 0$ 
4: for  $j = 1$  to  $A.length$  do
5:      $C[A[j]] = C[A[j]] + 1$ 
6: for  $i = 1$  to  $k$  do
7:      $C[i] = C[i] + C[i - 1]$ 
8: for  $j = A.length$  downto 1 do
9:      $B[C[A[j]]] = A[j]$ 
10:     $C[A[j]] = C[A[j]] - 1$ 
```

Time

$\Theta(k)$

$\Theta(n)$

$\Theta(k)$

$\Theta(n)$

- Runtime is $\Theta(n + k)$
 - Depends on two input parameters instead of just the problem size n .
 - This is $O(n)$ if $k = O(n)$.

► Stability

- CountingSort is stable: numbers with the same value appear in the output in **the same order as** they do **in the input** array.
 - The order of equal elements is preserved.
 - This property is relevant when **satellite data** (e.g. Java objects) is attached to keys being sorted.
 - We may think of the original order being used to break ties between elements with equal keys.
 - Works well for sorting emails according to (1) read/unread and (2) date.

► Radix Sort

- Stability helps for sorting numbers digit by digit (or English words letter by letter).
- Assume that each array element has d digits (from lowest significance to highest significance)

RADIXSORT(A, d)

1: **for** $i = 1$ to d **do**

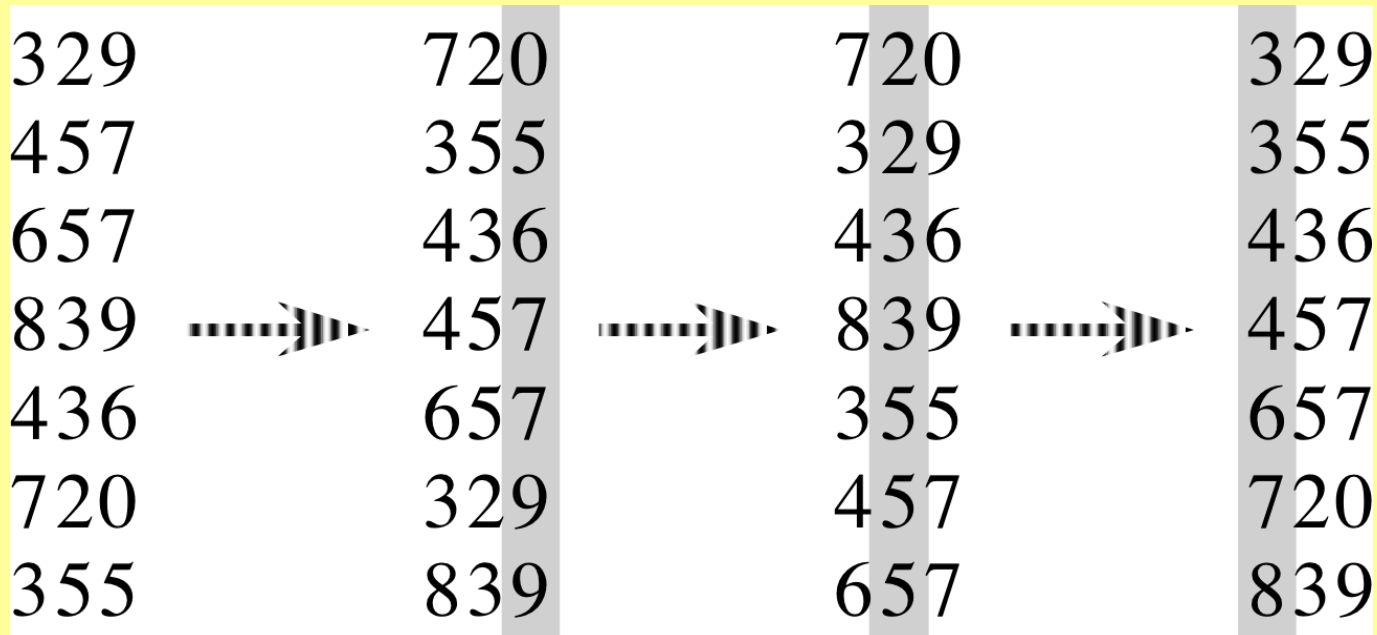
2: use a stable sort to sort array A on digit i

► Radix Sort: Example

RADIXSORT(A, d)

1: **for** $i = 1$ to d **do**

2: use a stable sort to sort array A on digit i



Correctness follows from stability and induction on columns.

► Radix Sort: Runtime

- Given n d -digit numbers in which each digit can take up to k possible values, RadixSort using CountingSort sorts these numbers in time $\Theta(d(n + k))$.
 - This is just the runtime of running CountingSort d times.

► Summary

- **Complexity Theory** gives limits to the efficiency of algorithms.
 - How (not) to prove lower bounds for all algorithms.
- All comparison sorts need time $\Omega(n \log n)$ in the worst case.
 - Decision trees capture the behaviour of every comparison sort.
- **CountingSort** sorts n numbers in a bounded range $\{0, \dots, k\}$ in time $\Theta(n + k)$.
- **RadixSort** uses a **stable sorting algorithm** to sort digit by digit.
 - **Stability** preserves the order of equal elements.
 - The time for sorting d -digit numbers is $\Theta(d(n + k))$.
 - This is $\Theta(n)$ when $d = O(1)$ and $k = O(n)$.