

# Parallel Computing with GPUs: Visual Studio Guide for CUDA

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



## Compiling a CUDA program

- ❑ CUDA C Code is compiled using **nvcc** e.g.
- ❑ Will compile host AND device code to produce an executable

```
nvcc -o example example.cu
```

- ❑ We will be using Visual Studio to build our CUDA code so we will not need to compile at the command line (unless you are running on ShARC)



## Creating a CUDA Project

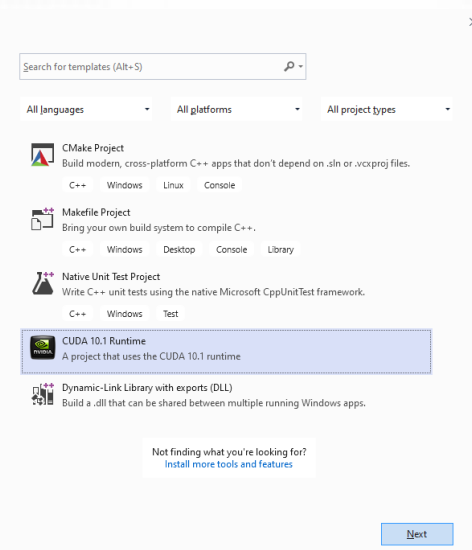
### ❑ Create New CUDA Project

- ❑ Select **NVIDIA -> CUDA 11.1**
- ❑ This will create a project with a default kernels.cu file containing a basic vector addition example

Create a new project

Recent project templates

- ❑ Empty Project C++
- ❑ CUDA 10.1 Runtime




Preferred Method!



## Adding a CUDA source file

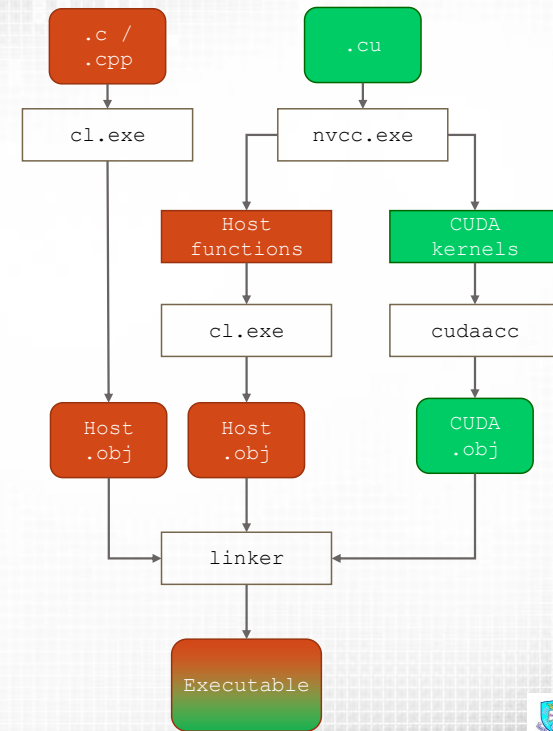
- ❑ Alternatively add a CUDA source file to an existing application
- ❑ If you do this you must modify the project properties to include CUDA build customisations

❑ [http://developer.download.nvidia.com/compute/cuda/6\\_5/rel/docs/CUDA\\_Getting\\_Started\\_Windows.pdf](http://developer.download.nvidia.com/compute/cuda/6_5/rel/docs/CUDA_Getting_Started_Windows.pdf) (section 3.4)



## Compilation

- ❑ CUDA source file (\* .cu) are compiled by `nvcc`
- ❑ An existing `cuda.rules` file creates property page for CUDA source files
  - ❑ Configures `nvcc` in the same way as configuring the C compiler
  - ❑ Options such as optimisation and include directories can be inherited from project defaults
- ❑ C and C++ files are compiled with `cl` (MSVCC compiler)



## Device Versions

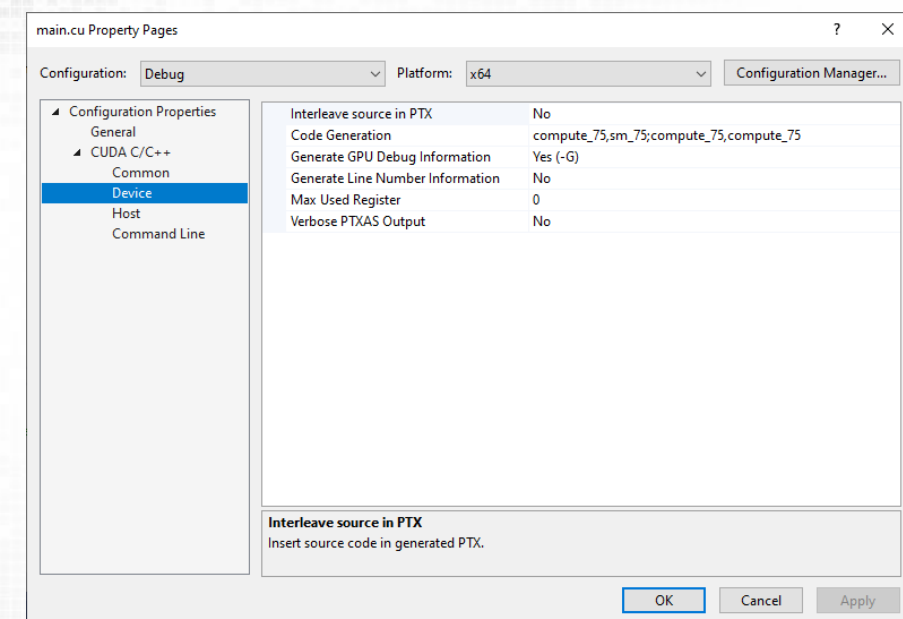
- ❑ Different generations of NVIDIA hardware have different compatibility
  - ❑ In the last lecture we saw product families and chip variants
  - ❑ These are classified by **CUDA compute versions**
- ❑ Compilation normally builds for CUDA compute version 35
  - ❑ See Project Properties, *CUDA C/C++ Device->Code Generation*
  - ❑ Default value is "compute\_35, sm\_35"
  - ❑ Any hardware with greater than the compiled compute version can execute the code (backwards compatibility)
- ❑ You can build for multiple versions using separator
  - ❑ E.g. "compute\_35, sm\_35; compute\_60, sm\_60"
  - ❑ This will increase build time and execution file size
  - ❑ Runtime will select the best version for your hardware

[https://en.wikipedia.org/wiki/CUDA#Supported\\_GPUs](https://en.wikipedia.org/wiki/CUDA#Supported_GPUs)

## Device Versions of Available GPUs


- ❑ Diamond Machines
  - ❑ Pascal Architecture
    - ❑ compute\_60, sm\_60;

## CUDA Properties



## Debugging

- ❑ NSIGHT is a GPU debugger for debugging GPU kernel code
  - ❑ It does not debug breakpoints in host code
- ❑ To launch select insert a breakpoint and select NSIGHT-> Start CUDA Debugging
  - ❑ You must be in the debug build configuration.
  - ❑ When stepping all warps except the debugger focus will be paused
- ❑ Use conditional breakpoints to focus on specific threads
  - ❑ Right click on break point and select Condition



## Error Checking

- ❑ `cudaError_t`: enumerator for runtime errors
  - ❑ Can be converted to an error string (`const char *`) using `cudaGetErrorString(cudaError_t)`
- ❑ Many host functions (e.g. `cudaMalloc`, `cudaMemcpy`) return a `cudaError_t` which can be used to handle errors gracefully

```
cudaError_t cudaStatus;  
  
cudaStatus = cudaMemcpy(dev_a, a, size * sizeof(int), cudaMemcpyHostToDevice);  
if (cudaStatus != cudaSuccess) {  
    //handle error  
}
```

- ❑ Kernels do not return an error but if one is raised it can be queried using the `cudaGetLastError()` function

```
addKernel<<<1, size>>>>(dev_c, dev_a, dev_b);  
cudaStatus = cudaGetLastError();
```

