


Parallel Computing with GPUs

Advanced OpenMP Part 1 – Parallel Reduction



Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

- ❑ Reduction
 - ❑ Perform a parallel reduction using the reduction clause
 - ❑ Recognise the limitations of the reduction functionality

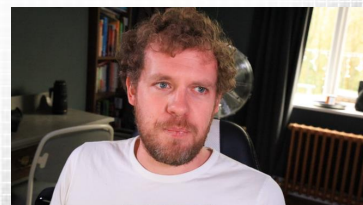


- ❑ What do we need to look out for when considering applying OpenMP to this example?

```
void main(){
    int i;
    float vector[N];
    float sum;

    init_vector_values(vector);
    sum = 0;

    for (i = 0; i < N; i++){
        float v = some_func(vector[i]);
        sum += v;
    }
    printf("Sum of values is %f\n", sum);
}
```



Parallel Reduction

- ❑ A Reduction is the combination of local copies of a variable into a single copy
 - ❑ Consider a case where we want to sum the values of a function operating on a vector of values;

```
void main(){
    int i;
    float vector[N];
    float sum;

    init_vector_values(vector);
    sum = 0;

    for (i = 0; i < N; i++){
        float v = some_func(vector[i]);
        sum += v;
    }
    printf("Sum of values is %f\n", sum);
}
```

Candidate for
parallel
reduction...



Reduction clause

```
void main(){
    int i;
    float vector[N];
    float sum;

    init_vector_values(vector);
    sum = 0;

    #pragma omp parallel for reduction(+: sum);
    for (i = 0; i < N; i++){
        float v = some_func(vector[i]);
        sum += v;
    }
    printf("Sum of values is %f\n", sum);
}
```

Without reduction
we would need a
critical section to
update the shared
variable!



OpenMP Reduction

- ❑ Reduction is supported with the reduction clause which requires a reduction variable
 - ❑ E.g. #pragma omp parallel reduction(+: sum_variable) {...}
 - ❑ Reduction variable is implicitly private to other threads
- ❑ OpenMP implements this **in parallel** by;
 - ❑ Creating a local (private) copy of the (shared) reduction variable
 - ❑ Combining (merging) local copies of the variable at the end of the structured block
 - ❑ Saving the reduced value to the shared variable in the master thread.
- ❑ Reduction operators are +, -, *, &, |, && and ||
 - ❑ &: bitwise and
 - ❑ |: bitwise or
 - ❑ &&: logical and
 - ❑ ||: logical or



Summary

- ❑ Reduction
 - ❑ Perform a parallel reduction using the reduction clause
 - ❑ Recognise the limitations of the reduction functionality

❑ Next Lecture: Scheduling



Parallel Computing with GPUs

Advanced OpenMP Part 2 – Scheduling



Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

❑ Scheduling

- ❑ Compare and contrast different scheduling approaches to understand the benefits and limitations of each
- ❑ Identify how scheduling parameters may impact cache utilisation



Scheduling clause



- ❑ OpenMP by default uses static scheduling
 - ❑ Static: schedule is determined at compile time
 - ❑ `schedule(static)`
- ❑ In general: `schedule(type [, chunk size])`
 - ❑ `type=static`: Iterations assigned to threads before execution (preferably at compile time)
 - ❑ `type=dynamic`: iterations are assigned to threads as they become available
 - ❑ `type=guided`: iterations are assigned to threads as they become available (with reducing chunk size)
 - ❑ `type=auto`: compiler and runtime determine the schedule
 - ❑ `type=runtime`: schedule is determined at runtime by env variable

What would be a use case where static scheduling is a bad choice?



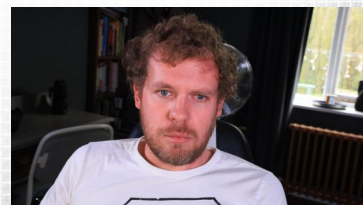
Static scheduling chunk size

❑ chunk size

- ❑ Refers to the amount of work assigned to each thread
- ❑ By default chunk size is to divide the work by the number of threads
 - ❑ Low overhead (no going back for more work)
 - ❑ Not good for uneven workloads
 - ❑ E.g. consider our last lectures Taylor series example (updated to use reduction)

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for reduction(-: result)  
for (n = 0; n < EXPANSION_STEPS; n++){  
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);  
    result -= r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

Recursive uneven workload



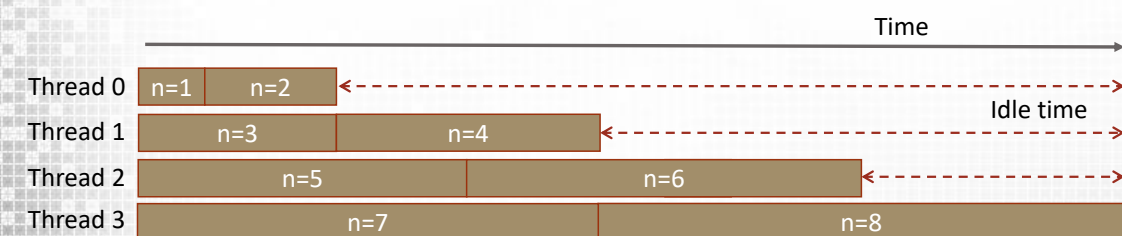
Scheduling Workload



```
long long int fac(int n)  
{  
    if (n == 0)  
        return 1;  
    else  
        return (n * fac(n - 1));  
}
```

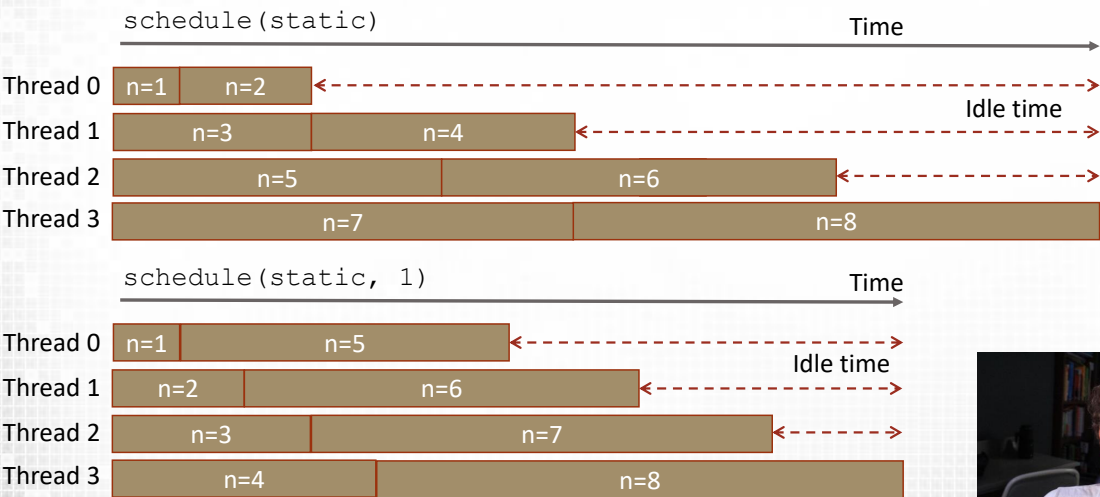
❑ Uneven workload amongst threads

- ❑ Increase in n leads to increased computation
- ❑ E.g. `EXPANSION_STEPS=8, num_threads(4), schedule(static)`



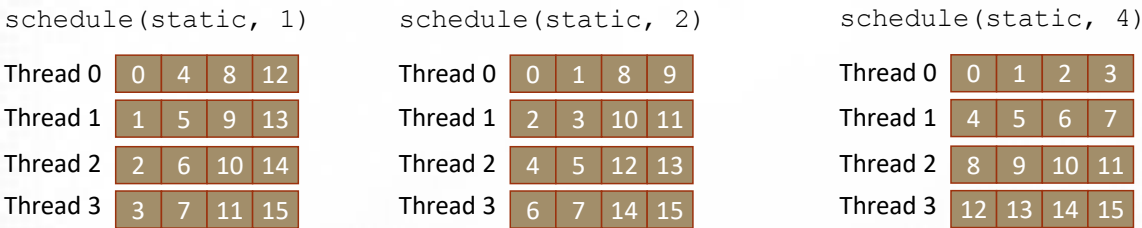
Cyclic Scheduling

- It would be better to partition the workload more evenly
 - E.g. Cyclic scheduling via chunk size



Cyclic Scheduling

```
#pragma omp for num_threads(4)
for (i = 0; i < 16; i++)
```



Default case

- Default chunk size is $n / \text{threads}$
 - where n is the number of iterations



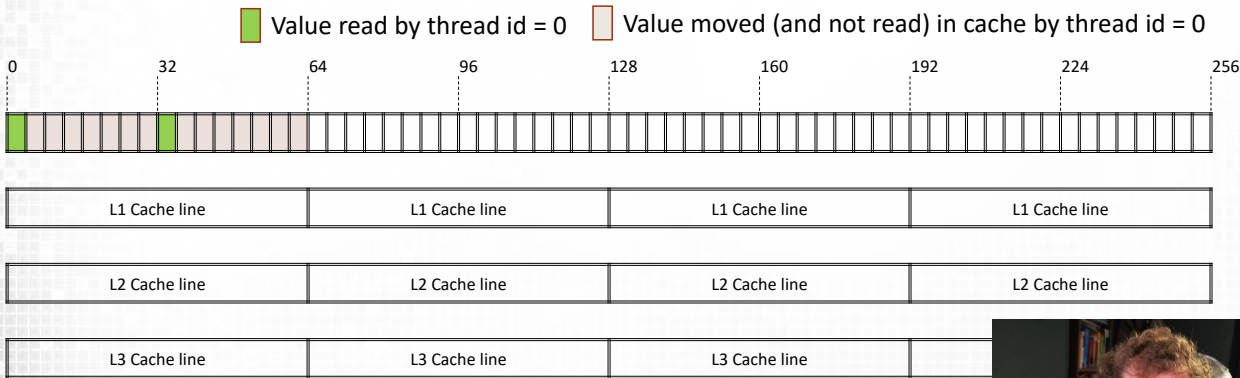
Dynamic and Guided Scheduling

- Dynamic
 - Iterations are broken down by chunk size
 - Threads request chunks of work from a runtime queue when they are free
 - Default chunk size is 1
- Guided
 - Chunks of the workload grow exponentially smaller
 - Threads request chunks of work from a runtime queue when they are free
 - Chunk size is the size which the workloads decrease to
 - with the exception of last chunk which may have remainder
- Both
 - Requesting work dynamically creates overhead
 - Not well suited if iterations are balanced
 - Overhead vs. imbalance: How do I decide which is best?
 - Benchmark all to find the best solution



Cache Efficiency

```
#pragma omp parallel for schedule(static,1) num_threads(8)
for (int i=0; i<64; i++) {
    something(array[i]);
}
```



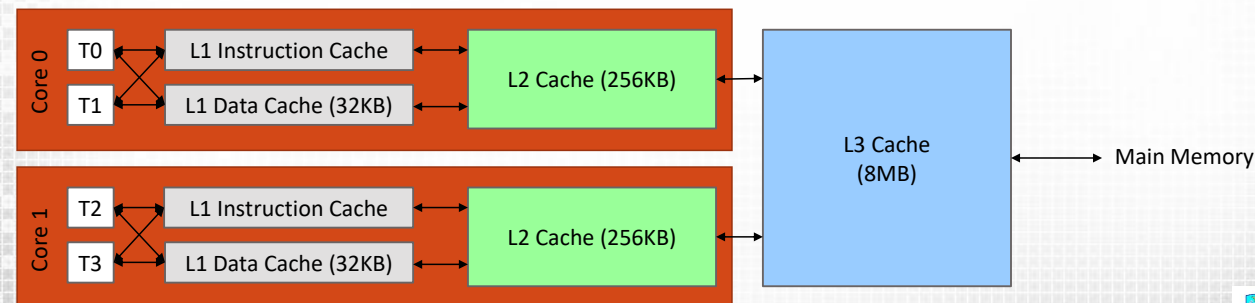
- Chunk size may effect cache utilisation



False Sharing

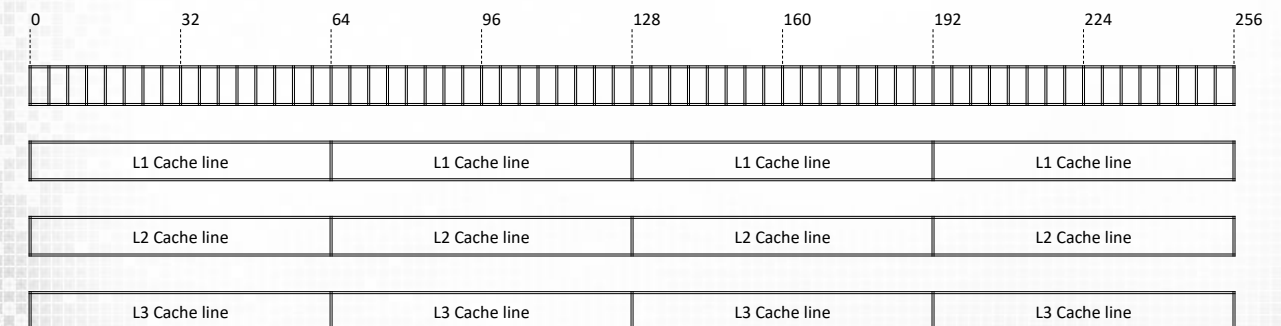
- ❑ Changing a single value causes a whole cache line to be invalid
 - ❑ Invalid lines must be re-cached
 - ❑ Chunk size case effect the amount of times a line is invalid

```
#pragma omp parallel for schedule(static,1)
for (int i=0; i<64; i++) {
    array[i]++;
}
```



False Sharing (worked example)

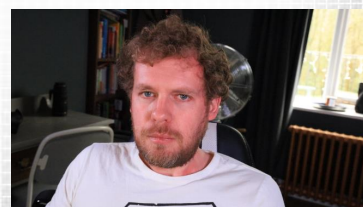
```
#pragma omp parallel for schedule(static,1) num_threads(8)
for (int i=0; i<64; i++) {
    array[i]++;
}
```



Summary

- ❑ Scheduling
 - ❑ Compare and contrast different scheduling approaches to understand the benefits and limitations of each
 - ❑ Identify how scheduling parameters may impact cache utilisation

❑ Next Lecture: Nesting Loops and OpenMP Summary



Parallel Computing with GPUs

Advanced OpenMP Part 3 – Nesting and Summary

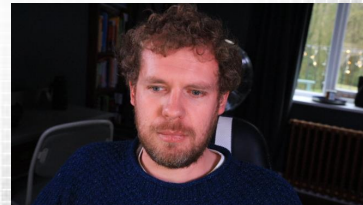


Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

- ❑ Nesting
 - ❑ Operate on nested loops using OpenMP
 - ❑ Compare the performance implications of different approaches for nesting
- ❑ Summary
 - ❑ Classify permitted use of the various OpenMP clauses



Nesting

- ❑ Consider the following example...
 - ❑ How should we parallelise this example?

```
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        printf("Hello World (Thread %d)\n", omp_get_thread_num());
    }
}
```

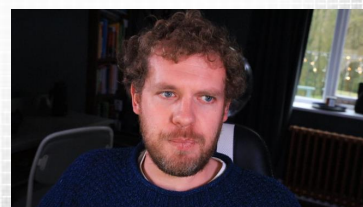


Nesting

- ❑ Consider the following example...
 - ❑ How should we parallelise this example?

```
#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        printf("Hello World (Thread %d)\n", omp_get_thread_num());
    }
}
```

- ❑ What if $OUTER_LOOPS \ll \text{number of threads}$
 - ❑ E.g. $OUTER_LOOPS = 2$



Nesting

- ❑ We can use parallel nesting
 - ❑ Nesting is turned off by default so we must use `omp_set_nested()`
 - ❑ When inner loop is met each outer thread creates a new team of threads
 - ❑ Allows us to expose higher levels of parallelism
 - ❑ Only useful when outer loop does not expose enough parallelism

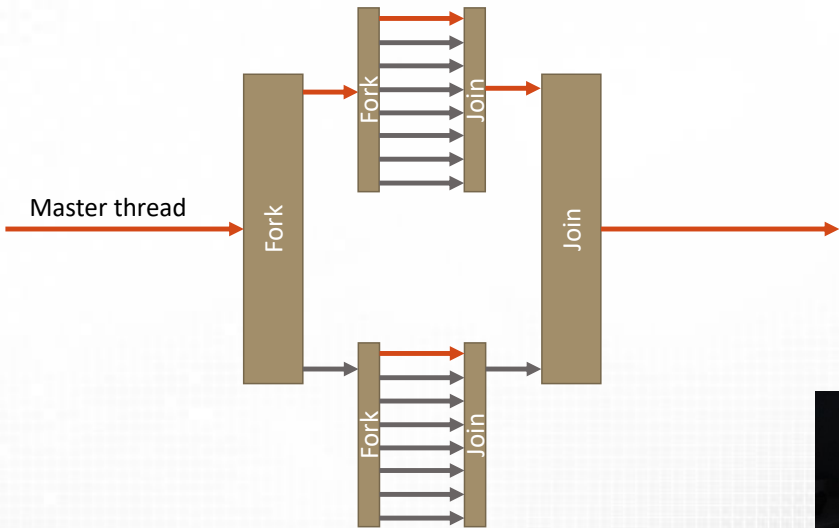
```
omp_set_nested(1);

#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS; i++){
    int outer_thread = omp_get_thread_num();
    #pragma omp parallel for
    for (j = 0; j < INNER_LOOPS; j++){
        int inner_thread = omp_get_thread_num();
        printf("Hello World (i T=%d j T=%d)\n", outer_thread, inner_thread);
    }
}
```

```
Hello World (i T=0 j T=0)
Hello World (i T=0 j T=1)
Hello World (i T=0 j T=3)
Hello World (i T=1 j T=2)
Hello World (i T=1 j T=1)
Hello World (i T=1 j T=0)
Hello World (i T=0 j T=2)
Hello World (i T=1 j T=3)
```


Nesting Fork and Join

- ❑ Every parallel directive creates a fork (new team)
- ❑ In this case `omp parallel` is used to fork a new parallel region



Collapse

- ❑ Only available in OpenMP 3.0 and later (not VS2017)
- ❑ Can automatically collapse multiple loops
- ❑ Loops must not have statements or expressions between them

```
#pragma omp parallel for collapse(2)
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        int thread = omp_get_thread_num();
        printf("Hello World (T=%d)\n", thread);
    }
}
```

Work around...

```
#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS* INNER_LOOPS; i++){
    int thread = omp_get_thread_num();
    printf("Hello World (T=%d)\n", thread);
}
```



Clauses usage summary

Clause	Directive: #pragma omp ...					
	parallel	for	sections	single	parallel for	parallel sections
if						
private						
shared						
default						
firstprivate						
lastprivate						
reduction						
schedule						
nowait						



Performance

- ❑ Remember ideas for general C performance
 - ❑ Have good data locality (good cache usage)
 - ❑ Combine loops where possible
- ❑ Additional performance criteria
 - ❑ Minimise the use of barriers
 - ❑ Use `nowait` but only if it is safe to do so!
 - ❑ Especially minimise critical sections
 - ❑ High overhead. Can you use reduction or atomics?
 - ❑ Benchmark to find best solution
 - ❑ Experimentally try out different scheduling approaches and chunk sizes



Summary

❑ Nesting

- ❑ Operate on nested loops using OpenMP
- ❑ Compare the performance implications of different approaches for nesting

❑ Summary

- ❑ Classify permitted use of the various OpenMP clauses



- ❑ Further Reading: <https://software.intel.com/en-us/articles/32-openmp-traps-for-c-developers>

