


Parallel Computing with GPUs

Memory Part 1 – Pointers

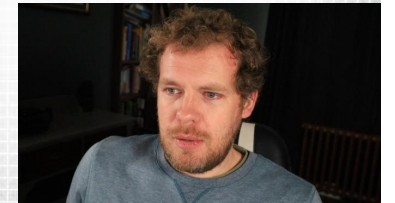


Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

- ❑ Pointers
 - ❑ Identify and use pointers and differentiate pointers from variables
- ❑ Pointers and arrays
 - ❑ Recognise the relationship between arrays and pointers
- ❑ Pointer arithmetic
 - ❑ Operate on pointers using simple arithmetic and predict how arithmetic operators make a pointers value change



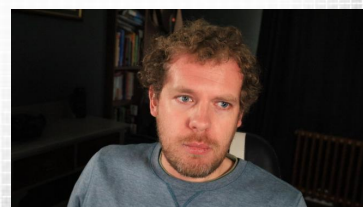
Pointers

- ❑ A pointer is a variable that contains the address of a variable
- ❑ Pointers and arrays are closely related
 - ❑ We have already seen some of the syntax with * and & operators
- ❑ The * operator can be used to define a pointer variable
- ❑ The operator & gives the address of a variable
 - ❑ Can not be applied to expressions or constants

```
#include <stdio.h>

void main()
{
    int a;
    int *p;

    a = 8;
    p = &a;
}
```



Pointer example

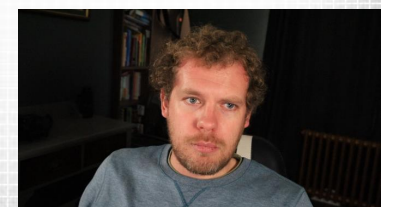
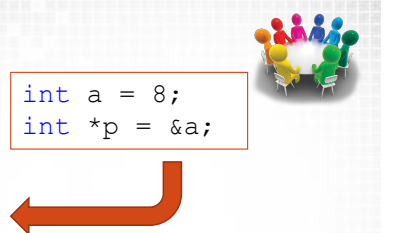
```
printf("a = %d, p = %d\n", a, p);
printf("a = %d, p = 0x%08X\n", a, p);
```

```
a = 8, p = 2750532
a = 8, p = 0x0045FCE0
```

❑ Same example using a char

```
char b;
char *p;
b = 8;
p = &b;
printf("sizeof(b) = %d, sizeof(p) = %d\n", sizeof(b), sizeof(p));
printf("b = %d, p = 0x%08X\n", b, p);
```

❑ What is the size of p?



Pointer example

```
printf("a = %d, p = %d\n", a, p);  
printf("a = %d, p = 0x%08X\n", a, p);
```

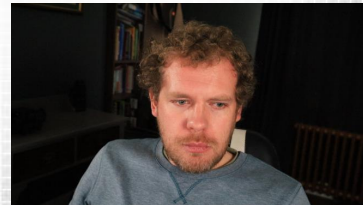
```
a = 8, p = 2750532  
a = 8, p = 0x0045FCE0
```

Same example using a char

```
char b;  
char *p;  
b = 8;  
p = &b;  
printf("sizeof(b) = %d, sizeof(p) = %d\n", sizeof(b), sizeof(p));  
printf("b = %d, p = 0x%08X\n", b, p);
```

What is the size of p?

```
sizeof(b) = 1, sizeof(p) = 4  
b = 8, p = 0x003BF9A7
```



Pointers

Pointer size does not change regardless of what it points to

- The size of a pointer on a 32 bit machine is always 4 bytes
- The size of a pointer on a 64 bit machine is always 8 bytes

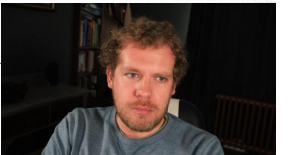
The operator * is the indirection operator and can be used to dereference a pointer

- I.e. it accesses the value that a pointer points to...

The macro NULL can be assigned to a pointer to give it a value 0

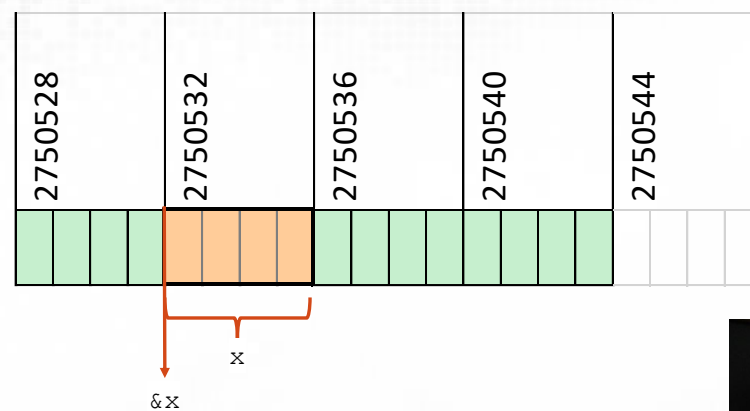
- This is useful in checking if a pointer has been assigned

```
int x = 1; int y = 0;  
int *p;  
p = &x; // p now points to x (value is address of x)  
y = *p; // y is now equal to the value of what p points to (i.e. x)  
x++;    // x is now 2 (y is still 1)  
(*p)++; // x is now 3 (y is still 1)  
p = NULL; // p is now 0
```

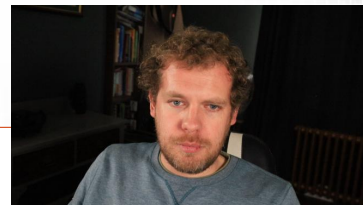


The University of Sheffield

Pointers



```
int x = 1; int y = 0;  
int *p;  
p = &x; // p now points to x (value is address of x)  
y = *p; // y is now equal to the value of what p points to (i.e. x)  
x++;    // x is now 2 (y is still 1)  
(*p)++; // x is now 3 (y is still 1)  
p = NULL; // p is now 0
```



The University of Sheffield

Pointers and arguments

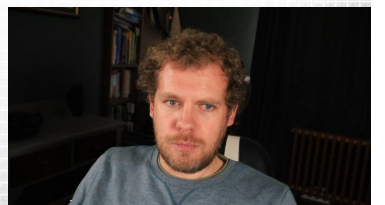
C passes function arguments by value

- They can therefore only be modified locally

```
void swap (int x, int y) {  
    int temp;  
    temp = x;  
    x = y;  
    y = temp;  
}
```

This is ineffective

- Local copies of x and y are exchanged and then discarded



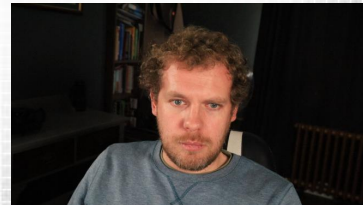
Pointers and arguments

- ❑ C passes function arguments by value
 - ❑ They can therefore only be modified locally

```
void swap (int *x, int *y){  
    int temp;  
    temp = *x;  
    *x = *y;  
    *y = temp;  
}
```

- ❑ This swaps the values which x and y point to
- ❑ Called by using the & operator

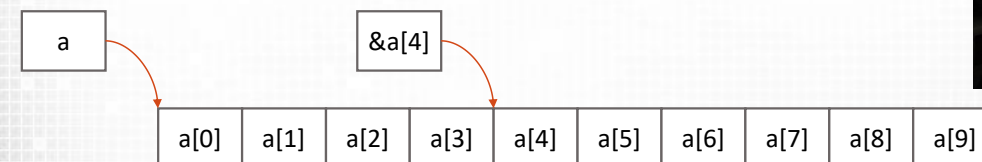
```
swap(&x, &y);
```



Pointers and Arrays

- ❑ In the last lecture we saw pointer being used for arrays
 - ❑ `char *name` is equivalent to `char name []`
- ❑ When we declare an array at compile time the variable **is a pointer** to the starting address of the array
 - ❑ E.g. `int a[10];`

```
int a[10] = {1,2,3,4,5,6,7,8,9,10};  
int *p;  
p = &a[4];  
printf("*p=%d, p[0]=%d\n", *p, p[0]);
```



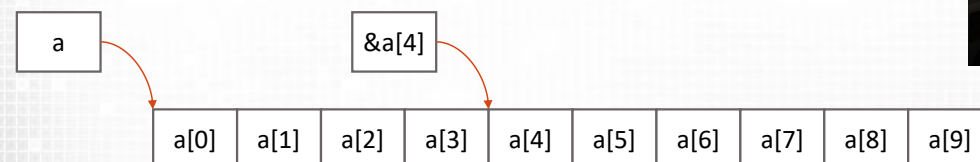
What is the output?




Pointers and Arrays

- ❑ In the last lecture we saw pointer being used for arrays
 - ❑ `char *name` is equivalent to `char name []`
- ❑ When we declare an array at compile time the variable **is a pointer** to the starting address of the array
 - ❑ E.g. `int a[10];`

```
int a[10] = {1,2,3,4,5,6,7,8,9,10};  
int *p;  
p = &a[4];  
printf("*p=%d, p[0]=%d\n", *p, p[0]);
```



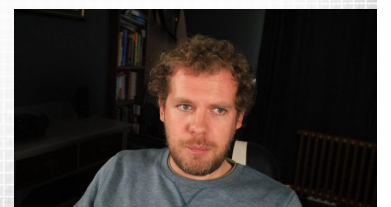
```
*p=5, p[0]=5
```



Pointer and Arrays

- ❑ There is however an important distinction between `char *name` and `char name []`
- ❑ Consider the following
 - ❑ The pointer may be modified
 - ❑ The array can only refer to the same storage

```
char a[] = "hello world 1";  
char *b = "hello world 2";  
char *temp;  
temp = b;  
b = a;  
a = temp; //ERROR
```

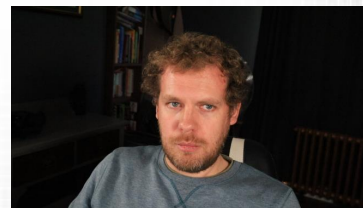
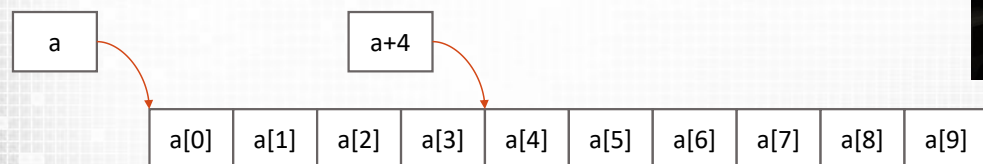


Pointer arithmetic



- ❑ Pointer can be manipulated like any other value
 - ❑ `p++`: advances the pointer the next element
 - ❑ Pointer arithmetic must not go beyond the bounds of an array
- ❑ Incrementing a pointer increments the memory location depending on the pointer type
 - ❑ An single integer *pointer* will increment 4 bytes to the next integer

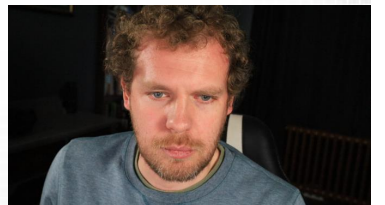
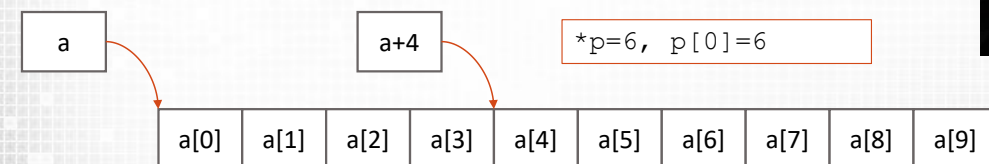
```
int a[10] = {10,9,8,7,6,5,4,3,2,1};
int *p = a;
p+=4;
printf("*p=%d, p[0]=%d\n", *p, p[0]);
```



Pointer arithmetic

- ❑ Pointer can be manipulated like any other value
 - ❑ `p++`: advances the pointer the next element
 - ❑ Pointer arithmetic must not go beyond the bounds of an array
- ❑ Incrementing a pointer increments the memory location depending on the pointer type
 - ❑ An single integer *pointer* will increment 4 bytes to the next integer


```
int a[10] = {10,9,8,7,6,5,4,3,2,1};
int *p = a;
p+=4;
printf("*p=%d, p[0]=%d\n", *p, p[0]);
```



This Lecture (learning objectives)

- ❑ Pointers
 - ❑ Identify and use pointers and differentiate pointers from variables
- ❑ Pointers and arrays
 - ❑ Recognise the relationship between arrays and pointers
- ❑ Pointer arithmetic
 - ❑ Operate on pointers using simple arithmetic and predict how arithmetic operators make a pointers value change

❑ Next Lecture: Advanced use of Pointers



Parallel Computing with GPUs

Memory Part 2 – Advanced use of Pointers



Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

- ❑Advanced use of Pointers
 - ❑Identify a general purpose pointer
 - ❑Determine the endianness of a computing system
 - ❑Interpret advanced pointer declarations
 - ❑Recognise function pointers and determine where they may be used



General Purpose Pointer

- ❑A General purpose pointer can be defined using `void` type
 - ❑A void type can not be dereferenced
 - ❑Carefull: Arithmetic on a void pointer will increment/decrement by 1 byte
 - ❑Even if it points to a 4 byte data type (e.g. int)

```
void *p;
char c;
int i;
float f;
p = &c; // ptr has address of character data
p = &i; // ptr has address of integer data
p = &f; // ptr has address of float data
```




Endianness

- ❑X86 uses little endian format
 - ❑Memory is stored from least significant byte stored at the **lowest** memory


```
unsigned int a = 0xDEADBEEF;
char* p;
p = (char*)&a; //Note explicit cast
printf("0x%08X, 0x%08X, 0x%08X, 0x%08X\n", p, p+1, p+2, p+3);
printf("0x%02X, 0x%02X, 0x%02X, 0x%02X\n", *p, *(p+1), *(p+2), *(p+3));
```

0x00400000, 0x00400001, 0x00400002, 0x00400003
0xEF, 0xBE, 0xAD, 0xDE



Endianness

```
int a[] = {0x12345678, 0x00000000, 0xDEADBEEF, 0xFFFFFFFF};
```



Endianness is very stange without an example
ssenaidsnE si yrev egnats tuohtiw na elpmaxe



Pointers to pointers

- ❑ Consider the following
 - ❑ `int a[10][20]`
 - ❑ `int *b[10]`
- ❑ `a` is a two-dimensional array
 - ❑ 200 int sized locations are reserved in memory
- ❑ `b` is single dimensional array of pointers
 - ❑ 10 pointers to integers are reserved
 - ❑ `B[?]` must be initialised (or allocated later in this lecture)
 - ❑ The pointers in `b` may be initialised to arrays of different length

```
char names[][10] = {"Paul", "Bob", "Emma", "Jim", "Kathryn"};
char *p_names[] = {"Paul", "Bob", "Emma", "Jim", "Kathryn"};
```

Which of the above is better?



Function Pointers

- ❑ It is possible to define pointers to functions
 - ❑ Functions are however **not** variables
- ```
int (*f_p)(int, int);
```
- ❑ `f_p` is a pointer to a function taking two integer arguments and returning an integer.
    - ❑ If `f` is a function then `&f` is a pointer to a function
    - ❑ Just in the same way that if `a` is an integer then `&a` is a pointer to an integer

```
int add(int a, int b);
int sub(int a, int b);

void main()
{
 int (*f_p)(int, int);
 f_p = &add;
 return;
}
```



## Using function pointers



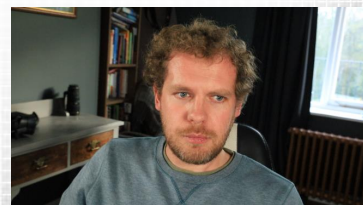
- ❑ Treat the function pointer like it is the function you want to call.
  - ❑ There is no need to dereference (`*f_p`) but you may if you wish

```
f_p = &add;
printf("add = %d\n", f_p(10, 4));
f_p = ⊂
printf("sub = %d\n", f_p(10, 4));
```

```
add = 14
sub = 6
```

- ❑ Care is needed with parenthesis
  - ❑ What is `f`?
  - ❑ What is `g`?

```
int *f();
int (*g)();
```



## Using function pointers

- ❑ Treat the function pointer like it is the function you want to call.
  - ❑ There is no need to dereference (`*f_p`) but you may if you wish

```
f_p = &add;
printf("add = %d\n", f_p(10, 4));
f_p = ⊂
printf("sub = %d\n", f_p(10, 4));
```

```
add = 14
sub = 6
```

- ❑ Care is needed with parenthesis
  - ❑ What is `f`? **function returning pointer to int**
  - ❑ What is `g`? **pointer to a function returning int**

```
int *f();
int (*g)();
```





## const pointers



- ❑ Remember the definition of `const`?
  - ❑ Not unintentionally modifiable
- ❑ What then is the meaning of the following?

```
char * const p;
```

```
const char * p;
```

```
char const * const p;
```



## const pointers

- ❑ Remember the definition of `const`?
  - ❑ Not unintentionally modifiable
  - ❑ Read from right to left

<https://cdecl.org/> - C Gibberish to English

- ❑ What then is the meaning of the following?

```
char * const p;
```

The pointer is constant but the data pointed to is not  
i.e. declare p as const pointer to char

```
char const * p;
```

The pointed to data is constant but the pointer is not  
i.e. declare p as pointer to const char

```
char const * const p;
```

The pointer is constant and the data it points to is also constant  
i.e. declare p as const pointer to const char

= 

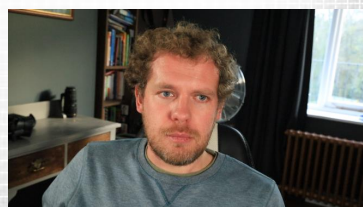
```
const char * p;
```



## Summary

- ❑ Advanced use of Pointers
  - ❑ Identify a general purpose pointer
  - ❑ Determine the endianness of a computing system
  - ❑ Interpret advanced pointer declarations
  - ❑ Recognise function pointers and determine where they may be used

- ❑ Next Lecture: Dynamically managed Memory

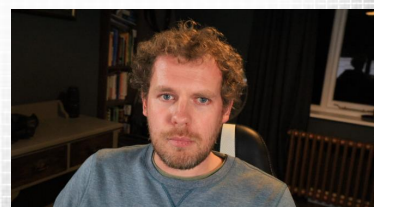


# Parallel Computing with GPUs

## Memory Part 3 – Dynamically managed memory



Dr Paul Richmond  
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



## This Lecture (learning objectives)

- ❑ Dynamically managed memory
  - ❑ Perform manual allocations and deletions of memory on the heap
  - ❑ Identify scenarios which may result in memory leaks
  - ❑ Operate on blocks of memory using library functions



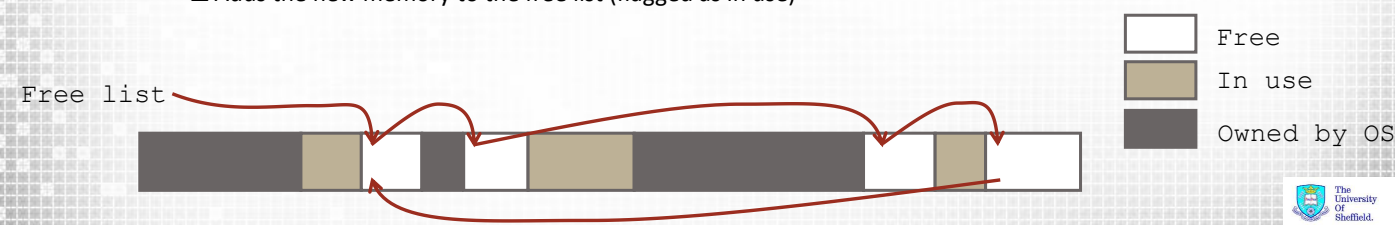
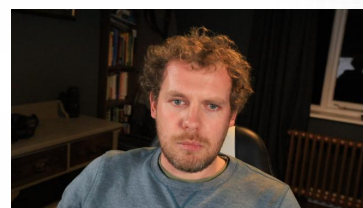
## Reminder: Heap vs. Stack

- ❑ Stack
  - ❑ Memory is managed for you
  - ❑ When a function declares a variable it is pushed onto the stack
  - ❑ When a function exists all variables on the stack are popped
  - ❑ Stack variables are therefore local
  - ❑ The stack has size limits
- ❑ Heap
  - ❑ You must manage memory
  - ❑ No size restrictions (except available memory)
  - ❑ Accessible by any function



## Dynamically allocated memory

- ❑ What if we can't specify an array size at compile time (static allocation)
  - ❑ The size might not be known until runtime
- ❑ We can use the `malloc` system function to get a block of memory on the heap.
  - ❑ `malloc` keeps a list of free blocks of memory on the heap
  - ❑ `malloc` returns the first free block which is big enough "e.g. first fit"
  - ❑ If a block is too big it is split
    - ❑ Part is returned to the user and the remainder added to the free list
  - ❑ If no suitable block is found `malloc` will request a larger block from the OS
    - ❑ Increases the size of the heap
    - ❑ Adds the new memory to the free list (flagged as in use)




## malloc

- ❑ `void *malloc(size_t size)`
  - ❑ Returns a pointer to void which must therefore be cast

```
#include <stdio.h>
#include <stdlib.h>

void main()
{
 int *a;
 a = (int*) malloc(sizeof(int) * 10);
}
```

- ❑ Use `sizeof` function to ensure correct number of bytes per element
- ❑ `a` can now be used as an array (as in the previous examples)
- ❑ Result of `malloc` will be implicitly cast (explicit cast is good practice)
  - ❑ Implicit cast generates a warning





# Memory leaks

- ❑ Consider the following
  - ❑ b is on the stack and is free'd on return
  - ❑ a points to an area of memory which is allocated
  - ❑ a then points to b, there is no pointer to the area of memory that was allocated

```
void main()
{
 int b[10] = {1,2,3,4,5,6,7,8,9,10};
 int *a;
 a = (int*) malloc(sizeof(int) * 10);
 a = b;

 return;
}
```

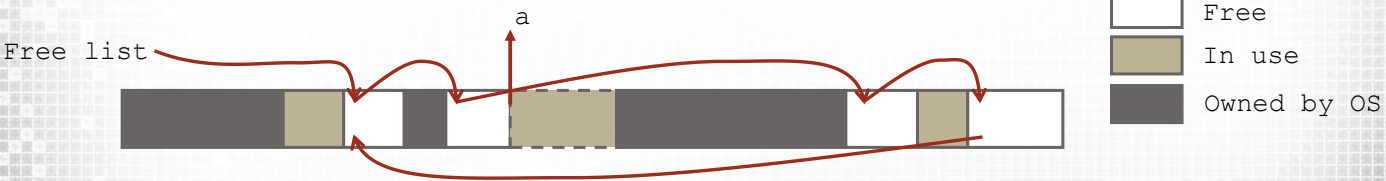
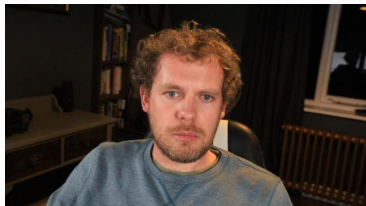
- ❑ This is known as a memory leak
  - ❑ Where we allocate memory we must also free it



# free

- ❑ The free function will add a previous used area of memory to the free list
  - ❑ If it is adjacent to another free block these will be coalesced into a larger block
- ❑ void free (void \*);

```
int *a = (int*) malloc(sizeof(int) * 10); //allocate
free(a); //free
```



# free

- ❑ The free function will add a previous used area of memory to the free list
  - ❑ If it is adjacent to another free block these will be coalesced into a larger block
- ❑ void free (void \*);

```
int *a = (int*) malloc(sizeof(int) * 10); //allocate
free(a); //free
```



# free

- ❑ The free function will add a previous used area of memory to the free list
  - ❑ If it is adjacent to another free block these will be coalesced into a larger block
- ❑ void free (void \*);

```
int *a = (int*) malloc(sizeof(int) * 10); //allocate
free(a); //free
```



## free

- ❑ The free function will add a previous used area of memory to the free list
  - ❑ If it is adjacent to another free block these will be coalesced into a larger block
- ❑ `void free (void *);`

```
int *a = (int*) malloc(sizeof(int) * 10); //allocate
free(a); //free
```



Free list



Free  
In use  
Owned by OS



## free

- ❑ The free function will add a previous used area of memory to the free list
  - ❑ If it is adjacent to another free block these will be coalesced into a larger block
- ❑ `void free (void *);`

```
int *a = (int*) malloc(sizeof(int) * 10); //allocate
free(a); //free
```



Free list



Free  
In use  
Owned by OS



## Memory operations

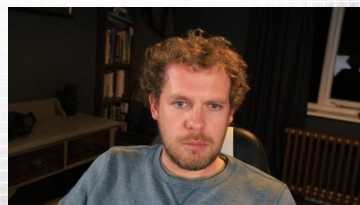
- ❑ Set a block of memory to char value
  - ❑ `void *memset(void *str, int c, size_t n)`
    - ❑ Can be used to set any memory to a value (e.g. 0)
    - ❑ Useful as allocated memory has undefined values

```
int *a;
int size = sizeof(int) * 10;
a = (int*) malloc(size);
memset(a, 0, size);
```

### Coping memory

- ❑ `void *memcpy(void *dest, const void *src, size_t n)`
  - ❑ Copies n bytes of memory from src to dst

```
int *a;
int b[] = {1,2,3,4,5,6,7,8,9,10};
int size = sizeof(int) * 10;
a = (int*) malloc(size);
memcpy(a, b, size);
```



## Summary

- ❑ Dynamically managed memory
  - ❑ Perform manual allocations and deletions of memory on the heap
  - ❑ Identify scenarios which may result in memory leaks
  - ❑ Operate on blocks of memory using library functions

❑ Next Lecture: Structures and binary files



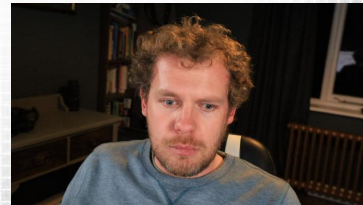


# Parallel Computing with GPUs

## Memory Part 4 – Structures and Binary Files



Dr Paul Richmond  
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



### This Lecture (learning objectives)

- ❑ Structures
  - ❑ Express a collection of variables as a structure and identify how to access member variables
- ❑ Binary Files
  - ❑ Apply functions to read and write to binary files

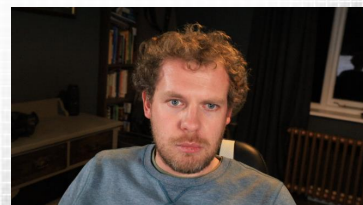


### Structures

- ❑ A structure is a collection of one or more variables
  - ❑ Variables may be of different types
  - ❑ Groups variables as a single unit under a single name
- ❑ A structure is not the same as a class (at least in C)
  - ❑ No functions
  - ❑ No private members
  - ❑ No inheritance
- ❑ Structures are defined using the `struct` keyword
  - ❑ Values can be assigned with an initialisation list or through structure member operator `'.'`

```
struct vec{
 int x;
 int y;
};

struct vec v_1 = {123, 456};
struct vec v_2;
v_2.x = 123;
v_2.y = 456;
```



### Features of structures

- ❑ As with everything, structures are passed by value
- ❑ Pointers to structures use a different member operator
  - ❑ `'->'` accesses member of a pointer to a struct
  - ❑ Alternatively dereference and use the standard operator `'.'`

```
struct vec make_vec(int x, int y){
 struct vec v = {x, y};
 return v;
}
```

- ❑ Declarations and definition can be combined

```
struct vec{
 int x;
 int y;
} v1 = {123, 456};
```



## Structure assignment



### ❑ Structures can be assigned

- ❑ Arithmetic operators not possible (e.g. `vec_2 += vec_1`)

```
struct vec vec_1 = {12, 34};
struct vec vec_2 = {56, 78};
vec_2 = vec_1;
```

### ❑ BUT No deep copies of pointer data

- ❑ E.g. if a person struct is declared with two char pointer members (forename and surname)

```
struct person paul, imposter;
paul.forename = (char *) malloc(5);
paul.surname = (char *) malloc(9);
strcpy(paul.forename, "Paul");
strcpy(paul.surname, "Richmond");
imposter = paul; // shallow copy
strcpy(imposter.forename, "John");
printf("Forename=%s, Surname=%s\n", paul.forename, paul.surname);
```

What is the Output?



## Structure assignment

### ❑ Structures can be assigned

- ❑ Arithmetic operators not possible (e.g. `vec_2 += vec_1`)

```
struct vec vec_1 = {12, 34};
struct vec vec_2 = {56, 78};
vec_2 = vec_1;
```

### ❑ BUT No deep copies of pointer data

- ❑ E.g. if a person struct is declared with two char pointer members (forename and surname)

```
struct person paul, imposter;
paul.forename = (char *) malloc(5);
paul.surname = (char *) malloc(9);
strcpy(paul.forename, "Paul");
strcpy(paul.surname, "Richmond");
imposter = paul; // shallow copy
strcpy(imposter.forename, "John");
printf("Forename=%s, Surname=%s\n", paul.forename, paul.surname);
```

Forename=John, Surname=Richmond



## Structure allocations

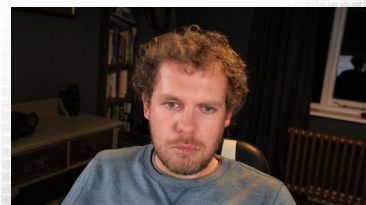
### ❑ Structures passed as arguments have member variables values copied

- ❑ If member is a pointer then pointer value copied not the thing that points to it (shown on last slide)
- ❑ Passing large structures by value can be quite inefficient

### ❑ Structures can be allocated and assigned to a pointer

- ❑ `sizeof` will return the combined size of all structure members
- ❑ Better to pass big structures as pointers

```
struct vec *p_vec;
p_vec = (struct vec *) malloc(sizeof(struct vec));
//...
free(p_vec);
```



## Type definitions

### ❑ The keyword `typedef` can be used to create 'alias' for data types

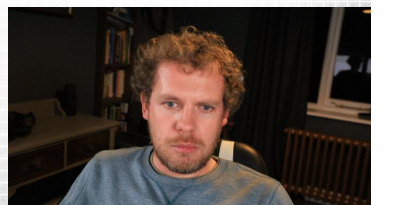
- ❑ Once defined a `typedef` can be used as a standard type

```
//declarations
typedef long long int int64;
typedef int int32;
typedef short int16;
typedef float vec3f [3];

//definitions
int32 a = 123;
vec3f vector = {1.0f, -1.0f, 0.0f};
```

### ❑ `typedef` is useful in simplifying the syntax of struct definitions

```
struct vec{
 int x;
 int y;
};
typedef struct vec vec;
vec p1 = {123, 456};
```





## Binary File Writing

- ❑ `size_t fwrite(const void *ptr, size_t size, size_t nmemb, FILE *stream)`
  - ❑ `size_t`: size of single object
  - ❑ `nmemb`: number of objects
  - ❑ Returns the number of objects written (if not equal to `nmemb` then error)

```
void write_points(FILE* f, point *points){
 fwrite(points, sizeof(point), sizeof(points) / sizeof(point), f);
}

void main(){
 point points[] = { 1, 2, 3, 4 };
 FILE *f = NULL;
 f = fopen("points.bin", "wb"); //write and binary flags
 write_points(f, points);
 fclose(f);
}
```



## Binary file reading

- ❑ `size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream)`

```
void read_points(FILE *f, point *points, unsigned int num_points){
 fread(points, sizeof(point), num_points, f);
}

void main(){
 point points[2];
 FILE *f = NULL;
 f = fopen("points.bin", "rb"); //read and binary flags
 read_points(f, points, 2);
 fclose(f);
}
```



## This Lecture (learning objectives)

- ❑ Structures
  - ❑ Express a collection of variables as a structure and identify how to access member variables
- ❑ Binary Files
  - ❑ Apply functions to read and write to binary files

