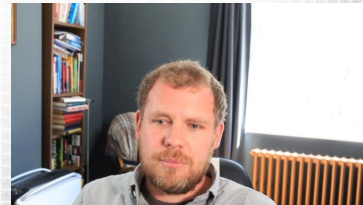


Parallel Computing with GPUs

Sorting and Libraries Part 1 – Sorting

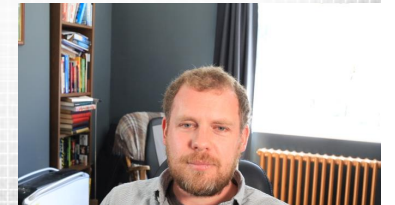


Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

- ❑ Sorting Networks
 - ❑ Demonstrate the use of a sorting network to achieve parallel sorting
 - ❑ Compare sorting networks with serial sorting approaches
- ❑ Merge and Bitonic Sort
 - ❑ Present the merge sort and considers its performance implications
 - ❑ Identify performance features of bitonic sorting



Serial Sorting Examples

❑ Insertion Sort

- ❑ Insert a new element into a sorted list.

❑ E.g. [1 6 3 4 2 5]

❑ [1] -> [1 6] -> [1 3 6] -> [1 3 4 6] -> [1 2 3 4 6] -> [1 2 3 4 5 6]

❑ Bubble Sort

- ❑ Exchange and Sweep to compare each pair of adjacent elements

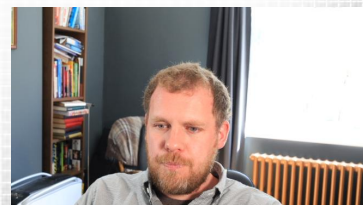
- ❑ $O(n^2)$ worst-case and average case, $O(n)$ best case.

❑ E.g. [1 6 3 4 2 5]

❑ [1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

❑ [1 3 2 4 5 6]

❑ [1 2 3 4 5 6]



Serial Sorting Examples

❑ Insertion Sort

- ❑ Insert a new element into a sorted list.

❑ E.g. [1 6 3 4 2 5]

❑ [1] -> [1 6] -> [1 3 6] -> [1 3 4 6] -> [1 2 3 4 6] -> [1 2 3 4 5 6]

❑ Bubble Sort

- ❑ Exchange and Sweep to compare each pair of adjacent elements

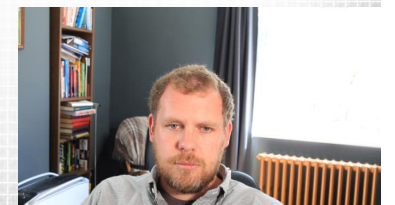
- ❑ $O(n^2)$ worst-case and average case, $O(n)$ best case.

❑ E.g. [1 6 3 4 2 5]

❑ [1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

❑ [1 3 2 4 5 6]

❑ [1 2 3 4 5 6]



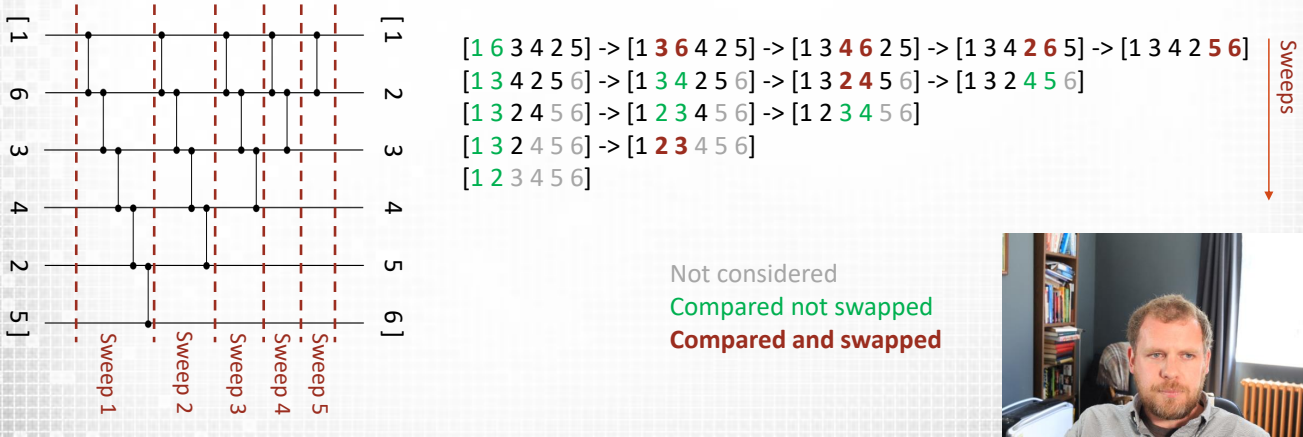
Classifying Sort Techniques/Implementations

- ❑ Data driven
 - ❑ Each step of the algorithm depends on the previous step version
 - ❑ Highly serial
- ❑ Data independent
 - ❑ The algorithms performs fixed steps and does not change its processing based on data
 - ❑ Well suited to parallel implementations
 - ❑ Can be expressed as a sorting network...



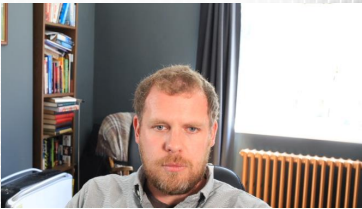
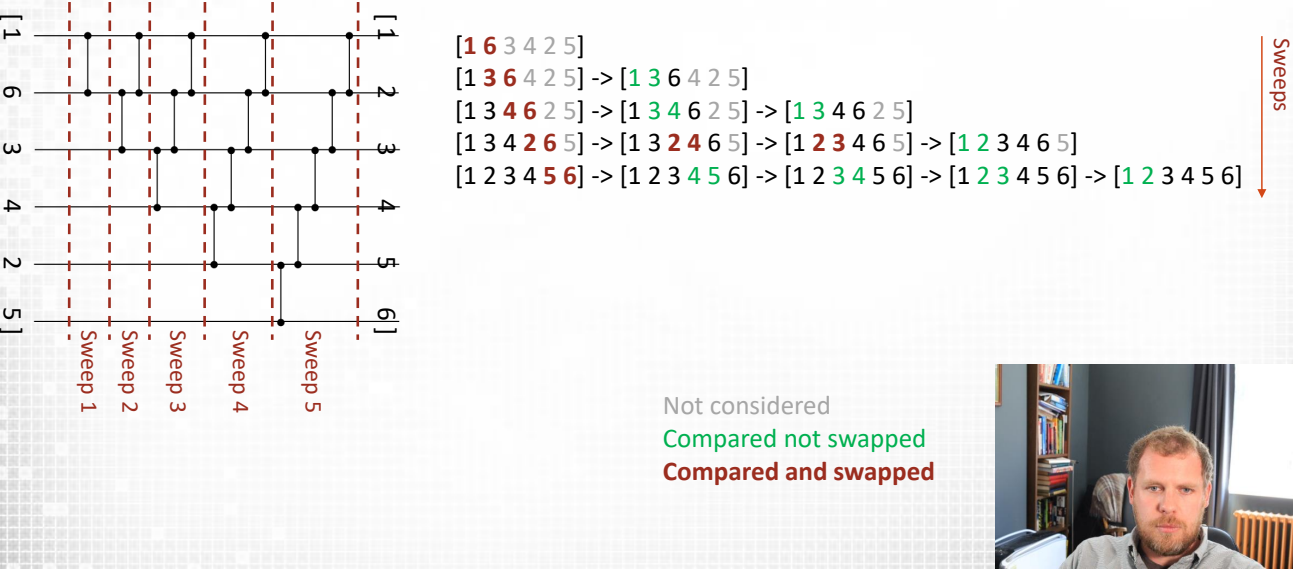
Sorting Networks

- ❑ A sorting network is a comparator network that sorts all input sequences
 - ❑ Following the same execution of stages
- ❑ Consider the previous Bubble Sort [1 6 3 4 2 5]

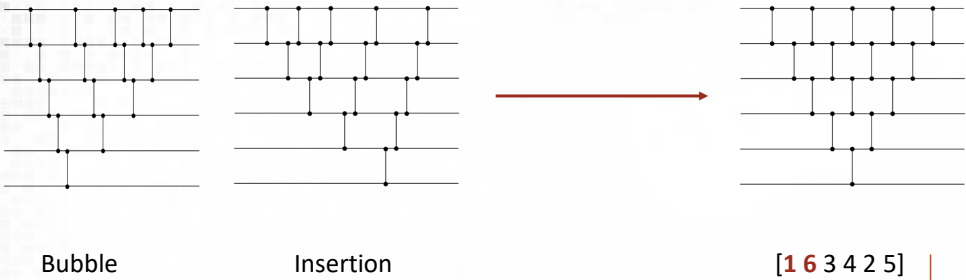


Sorting Networks

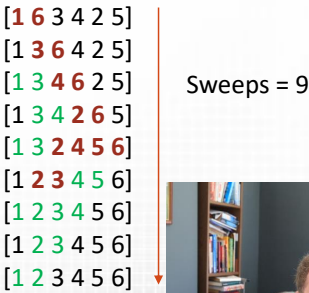
- ❑ And Insertion Sort...



Parallel Sorting Networks

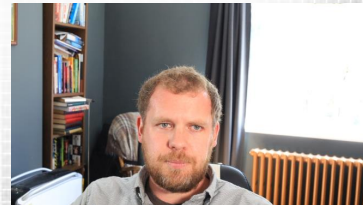
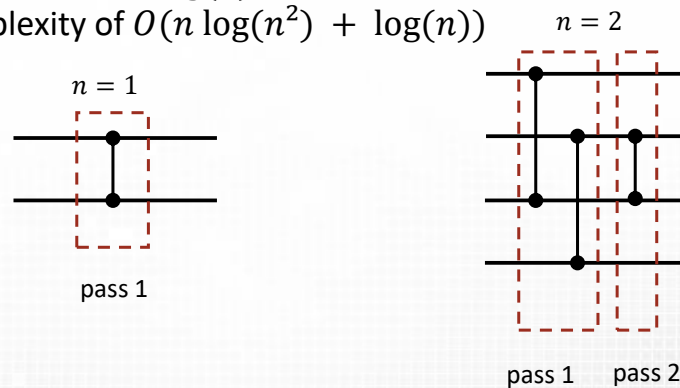


- ❑ Parallel Bubble and Insertion sorting network is still not very efficient
 - ❑ $2n - 3$ sweeps
 - ❑ $n(n - 1)/2$ comparisons - $O(n^2)$ complexity



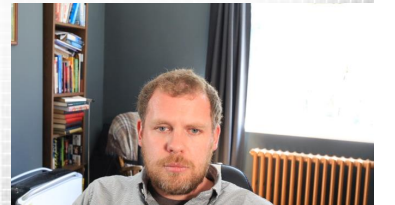
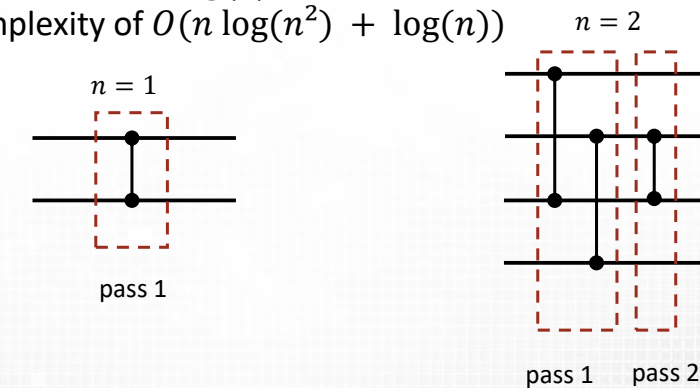
Merge Sort

- ❑ To reduce the $O(n^2)$ overhead we need a better sorting network
- ❑ The odd-even merge sort network (for power of 2 n)
 - ❑ Sort all odd and even keys separately and then merge m values of a stage
 - ❑ Merge a sorted sequence of elements on lines $\langle a_1, \dots, a_n \rangle$ with those on lines $\langle a_{n+1}, \dots, a_{2n} \rangle$
 - ❑ Each merge requires $\log(n)$ passes
 - ❑ Total complexity of $O(n \log(n^2) + \log(n))$

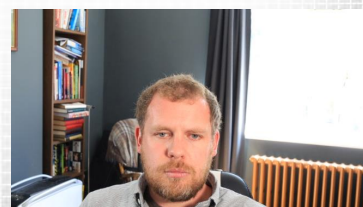
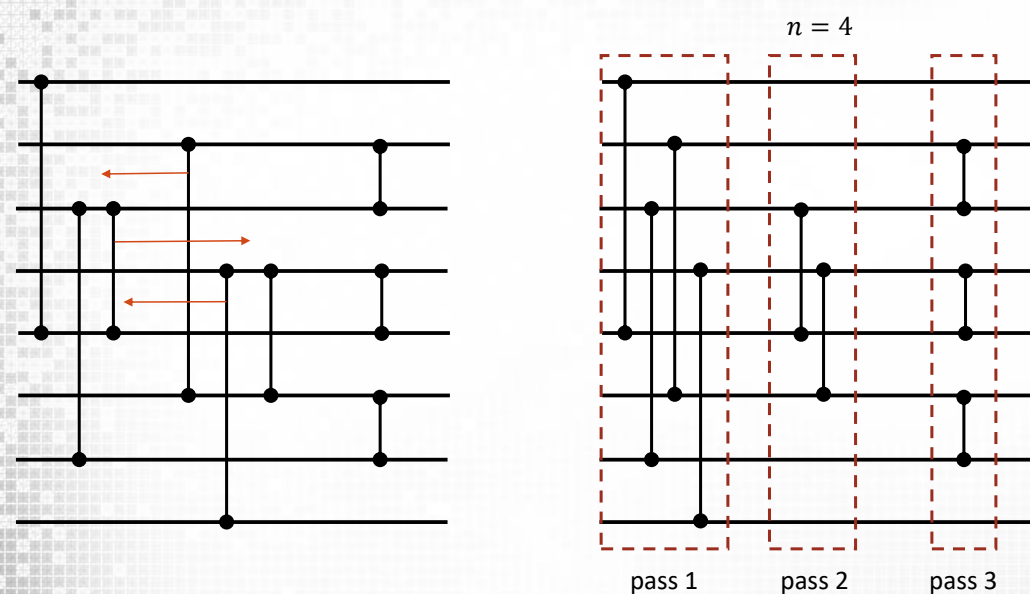


Merge Sort

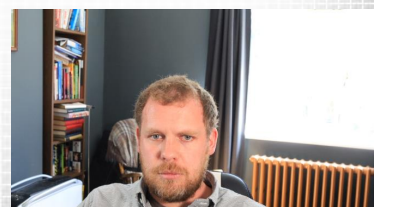
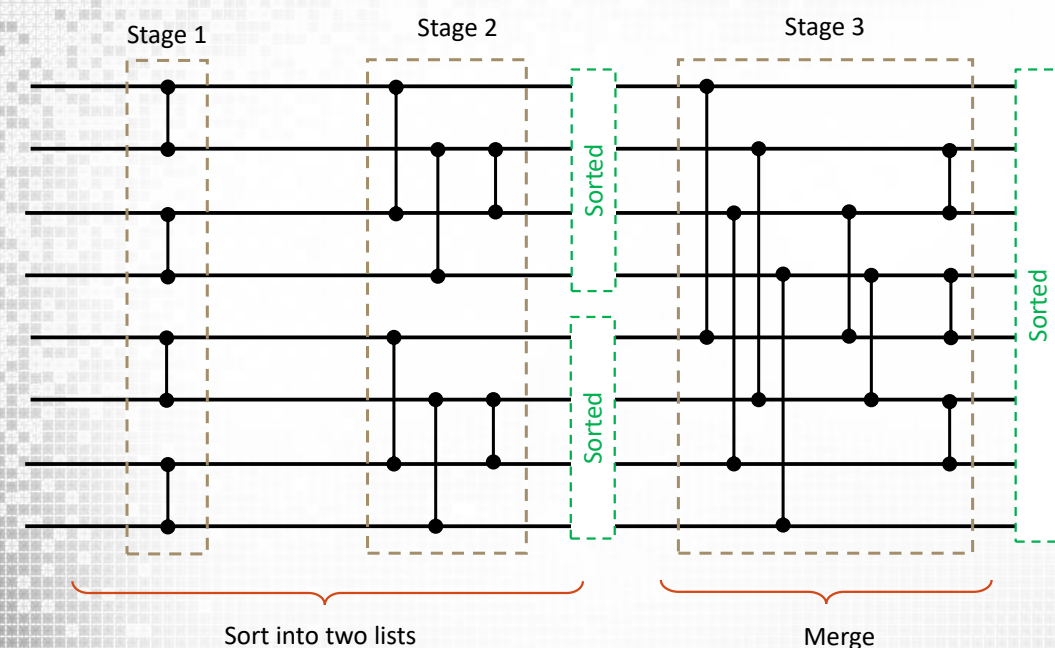
- ❑ To reduce the $O(n^2)$ overhead we need a better sorting network
- ❑ The odd-even merge sort network (for power of 2 n)
 - ❑ Sort all odd and even keys separately and then merge m values of a stage
 - ❑ Merge a sorted sequence of elements on lines $\langle a_1, \dots, a_n \rangle$ with those on lines $\langle a_{n+1}, \dots, a_{2n} \rangle$
 - ❑ Each merge requires $\log(n)$ passes
 - ❑ Total complexity of $O(n \log(n^2) + \log(n))$



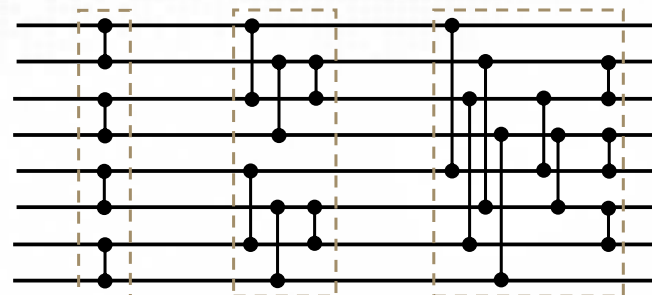
Merging of two sorted sequences (n=4)



Full Merge Sorting (n=4)



Full Merge Sorting (n=4) example



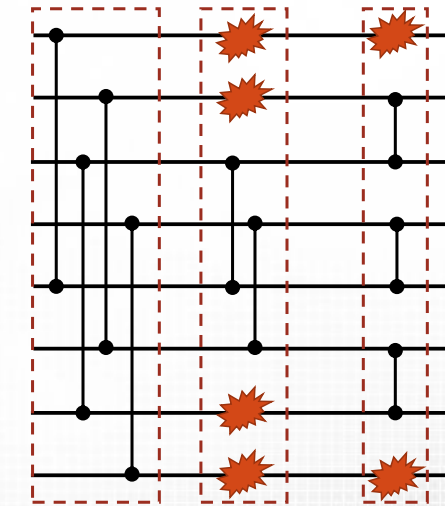
Input	Stage 1	Stage 2	Stage 3	Output
8	1	1	1	1
1	8	5	3	2
5	3	3	5	3
3	5	8	2	4
6	2	2	8	5
2	6	6	5	6
4	4	4	4	8
9	9	9	9	9



Limitations of Merge Sort

❑ What is potentially wrong with a merge sort GPU implementation?

- ❑ Irregular memory accesses
- ❑ Not all values are compared in each pass (uneven workload per thread)



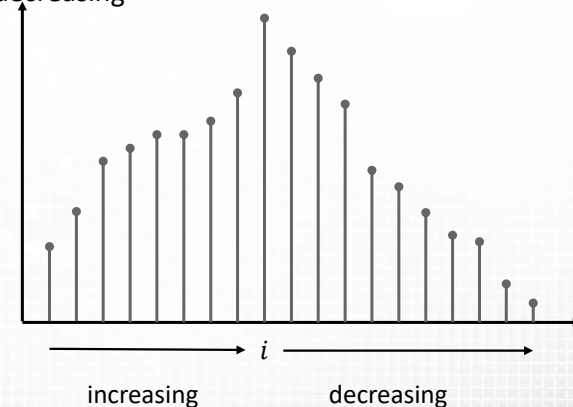
Solution: Bitonic Sort

❑ Bitonic sorting network

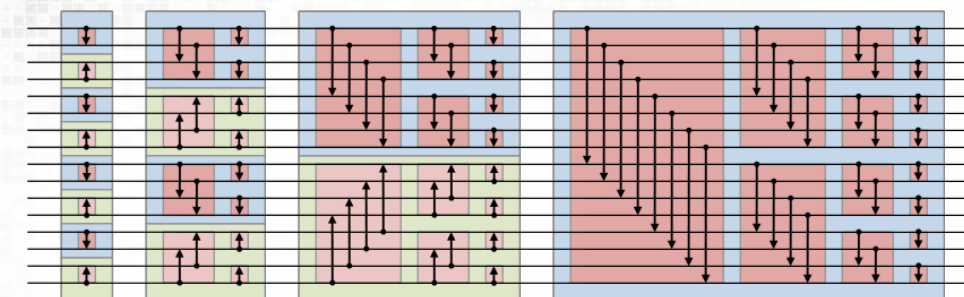
❑ Iterative splitting and merging of inputs into increasing large bionic sequences

❑ A sequence is bitonic if

❑ There is an i , such that $a_0 \dots, a_i$ is monotonically increasing and $a_i \dots, a_n$ is monotonically decreasing



Bitonic Sorting Network



❑ Sorting and Merging increasing large bionic sequences

❑ When $n = 2^k$ there are k levels with $\frac{n}{2}$ comparisons each

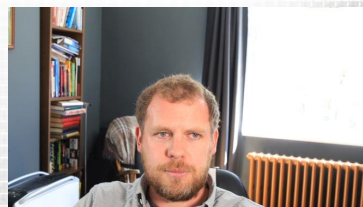
❑ GPU Implementation

- ❑ Regular access strides :-)
- ❑ Efficiently balanced workload :-)
- ❑ Requires multiple kernel launches to merge over $n > \text{block size}$

Summary

- ❑ Sorting Networks
 - ❑ Demonstrate the use of a sorting network to achieve parallel sorting
 - ❑ Compare sorting networks with serial sorting approaches
- ❑ Merge and Bitonic Sort
 - ❑ Present the merge sort and considers its performance implications
 - ❑ Identify performance features of bitonic sorting

❑ Next: Libraries and Thrust

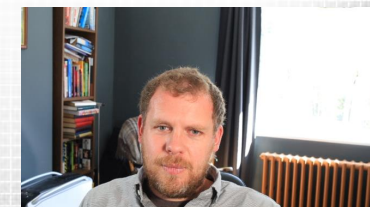


Parallel Computing with GPUs

Sorting and Libraries Part 2 – Libraries

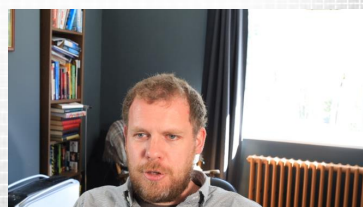


Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



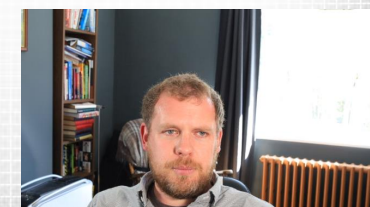
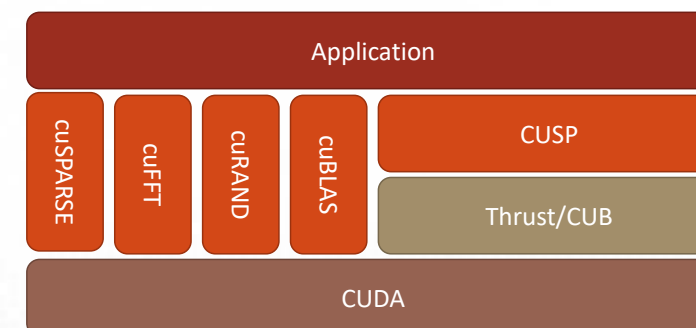
This Lecture (learning objectives)

- ❑ Libraries and Thrust
 - ❑ Describe the purpose of CUDA libraries
 - ❑ Demonstrate Thrust containers for data storage
 - ❑ Explain the relationship between raw pointers and Thrust iterators
 - ❑ Give example of Thrust algorithms



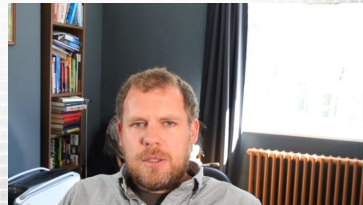
CUDA libraries

- ❑ Abstract CUDA model away from programmer
- ❑ Highly optimised implementations of common tools
 - ❑ Mainly focused on linear algebra



Thrust

- ❑ Template Library for CUDA
 - ❑ Implements many parallel primitives (scan, sort, reduction etc.)
 - ❑ Part of standard CUDA release
 - ❑ Level of Abstraction which hides kernels, mallocs and memcpy's
- ❑ Designed for C++ programmers
 - ❑ Similar in design and operation as the C++ Standard Template Library (STL)
 - ❑ Only a small amount of C++ required..



Thrust containers

- ❑ Thrust uses only high level *vector* containers
 - ❑ `host_vector`: on host
 - ❑ `device_vector`: on GPU
- ❑ Other STL containers include
 - ❑ `queue`
 - ❑ `list`
 - ❑ `tack`
 - ❑ `queue`
 - ❑ `priority_queue`
 - ❑ `set`
 - ❑ `multiset`
 - ❑ `map`
 - ❑ `multimap`
 - ❑ `bitset`
- ❑ STL containers can be used to initialise a Thrust vector

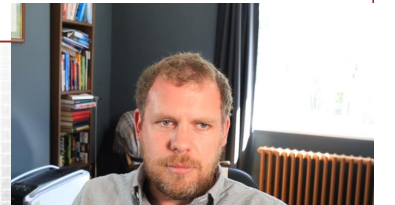
```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>

int main()
{
    //create a vector on the host
    thrust::host_vector<int> h_vec(10);

    //create a vector on the device
    thrust::device_vector<int> d_vec = h_vec;

    //device data manipulated directly from host
    for (int i = 0; i < 10; i++)
        d_vec[i] = i;

    //vector memory automatically released
    return 0;
}
```

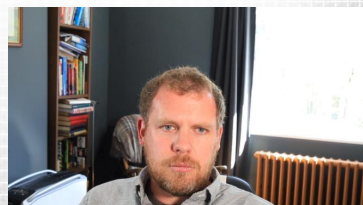


Thrust Iterators

- ❑ They point to regions of a vector
- ❑ Can be used like pointers
 - ❑ Explicit cast when dereferencing very important

```
thrust::device_vector<int>::iterator begin = d_vec.begin();
thrust::device_vector<int>::iterator end = d_vec.end();
printf("d_vec at begin=%d", (int)*begin);
begin++; //move on a single position
printf("d_vec at begin+=%d", (int)*begin);
*end = 88;
printf("d_vec at end=%d", (int)*end);
```

```
d_vec at begin=0
d_vec at begin+=1
d_vec at end=88
```



Thrust Iterators

- ❑ Can be converted to a raw pointer

```
int * d_ptr = thrust::raw_pointer_cast(begin);
int * d_ptr = thrust::raw_pointer_cast(begin[0]);

kernel<BLOCKS, TPB>(d_ptr);
```

- ❑ Raw pointers can be used in Thrust
 - ❑ BUT not exactly the same as a vector

```
int* d_ptr;
cudaMalloc((void**)&d_ptr, N);

thrust::device_ptr<int> d_vec = thrust::device_pointer_cast(d_ptr);
//or
thrust::device_ptr<int> d_vec = thrust::device_ptr<int>(d_ptr)
```


Thrust Algorithms

- ☐ Transformations
 - ☐ Application of a function to each element within the range of a vector
- ☐ Reduction
 - ☐ Reduction of a set of values to a single value using binary associative operator
 - ☐ Can also be used to count occurrences of a value
- ☐ Prefix Sum
 - ☐ Both inclusive and exclusive scans
- ☐ Sort
 - ☐ Can sort keys or key value pairs
- ☐ Binary Search
 - ☐ Position of a target value



Thrust Transformations

- ☐ Some examples of the many transformations

```
//copy a vector (or part of a vector) to another vector
thrust::copy(d_vec.begin(), d_vec.begin() + 10, d_vec_cpy.begin());

//fill a vector with a value
thrust::fill(d_vec.begin(), d_vec.begin() + 10, 0);

//rand is a predefined Thrust generator
thrust::generate(d_vec.begin(), d_vec.begin() + 10, rand);

// fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

//all occurrences of the value 1 are replaced with the value 10
thrust::replace(d_vec.begin(), d_vec.end(), 1, 10);
```



Thrust Algorithms

- ☐ Either in-place or to output vector

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);

//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

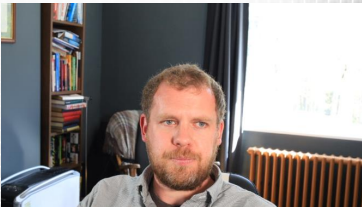
//inclusive scan to output vector
thrust::inclusive_scan(d_vec.begin(), d_vec.end(), d_vec_out.begin());

//inclusive scan in place
thrust::inclusive_scan(d_vec.begin(), d_vec.end(), d_vec.begin());

//generate random data (actually a transformation)
thrust::generate(d_vec.begin(), d_vec.end(), rand);

//sort in place
thrust::sort(d_vec.begin(), d_vec.end());

//sort data from a raw pointer (N is number of elements)
thrust::device_ptr<int> dt_ptr = thrust::device_pointer_cast(d_a_ptr);
thrust::sort(dt_ptr, dt_ptr+N);
```



Custom Transformations

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);

//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
d_vec = thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

//declare a custom operator
struct add_5{
    __host__ __device__ int operator()(int a){
        return a + 5;
    }
};

add_5 func;

//apply custom transformation
thrust::transform(d_vec.begin(), d_vec.end(), d_vec_out.begin(), func);

//d_vec is now {5, 6, 7, 8, 9, 10, 11, 12, 13, 14}
```



Thrust Fusion

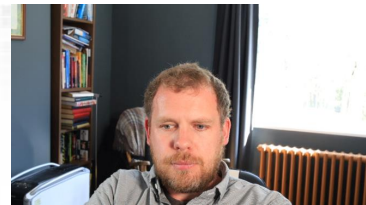
- ❑ For best performance it is necessary to fuse operations

```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;
```

```
//custom transformation to calculate absolute value
thrust::transform(d_vec.begin(), d_vec.end(), d_vec.begin(), func);
//apply reduction, maximum binary associate operator
int result = thrust::reduce(d_vec.begin(), d_vec.end(), 0, thrust::maximum<int>());
```

```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

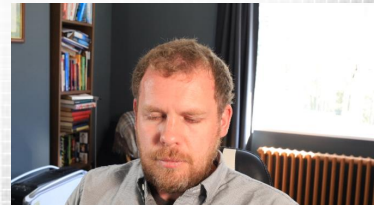
//apply transform reduction maximum binary associate operator
int result = thrust::transform_reduce(d_vec.begin(), d_vec.end(), func, 0, thrust::maximum<int>());
```



Summary

- ❑ Libraries and Thrust
 - ❑ Describe the purpose of CUDA libraries
 - ❑ Demonstrate Thrust containers for data storage
 - ❑ Explain the relationship between raw pointers and Thrust iterators
 - ❑ Give example of Thrust algorithms

- ❑ Next: Applications of Sort

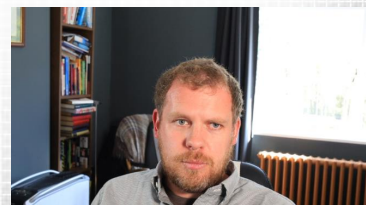


Parallel Computing with GPUs

Sorting and Libraries Part 3 – Applications of Sorting

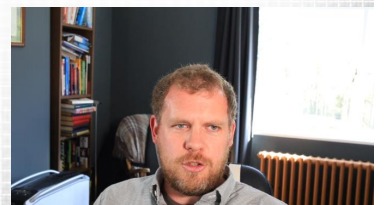


Dr Paul Richmond
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



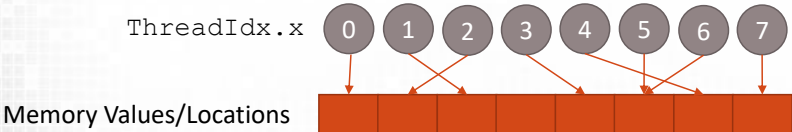
This Lecture (learning objectives)

- ❑ Applications of Sorting (binning)
 - ❑ Present the concept of spatial binning
 - ❑ Demonstrate the use of spatial binning for particle interactions

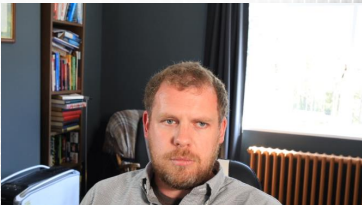


Sorting and parallel primitives

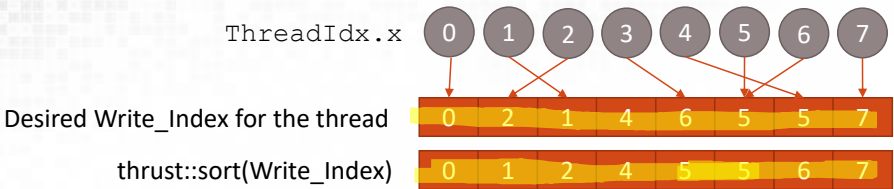
- Can be very useful for building data structures
 - We can use prefix sum for writing multiple values per element
- Remember Gather vs Scatter
 - What if our outputs are scattered to output
 - Very common in particle simulations etc.
 - Outputs might represent spatial bins



- Scatter operation
- Write to a number of locations
 - Random access write?
- How to read multiple values afterwards?



Binning and Sorting



Build a data structure

Unique write indices	0	1	2	3	4	5	6	7
Count(Write_Index)	1	1	1	0	1	2	1	1
thrust::inclusive_scan(count)	0	1	2	3	3	4	6	7

i.e. how many threads want to write to this index

- We can now read varying values from each bin
 - E.g. for location 5
 - inclusive_scan gives starting index of 4
 - Iterate from index 4 for a count of 2 to find all values of write_index 5



Particle interaction example

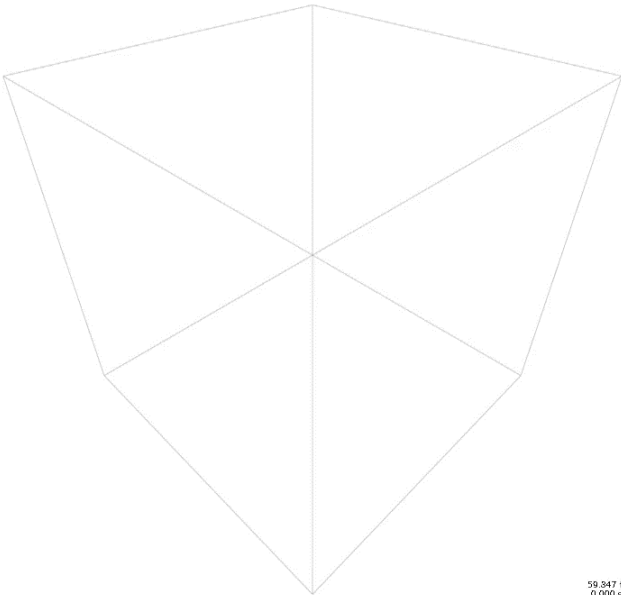
- As with previous slide use sorting
 - Divide the environment according to some interaction radius
 - Output particle key value pairs (keys are location determined through some hash function)
- Sort Keys
- Reorder particles based on key pairs
- Generate a partition boundary table
 - Histogram count and prefix sum
- Each particle needs to read all particles in its own location and any neighbouring location
 - Guarantees particle interactions within the interaction radius

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Partition	First agent	Last agent
0	0	0
1	1	2
2	3	4
3	5	6
4	7	7
5	8	8
6		
7		
8		
9		
10		
11		
12		
13		
14		
15		



Example: FLAME GPU



59.347 fps
0.000 eps
Step 1



Summary

- ❑ Applications of Sorting (binning)
 - ❑ Present the concept of spatial binning
 - ❑ Demonstrate the use of spatial binning for particle interactions

